

The OSU College of Engineering DGX System for Advanced GPU Computing



Oregon State
University
Mike Bailey

mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons
Attribution-NonCommercial-NoDerivatives 4.0
International License](#)

OSU's College of Engineering has *six* Nvidia DGX-2 systems

2

Each DGX server:

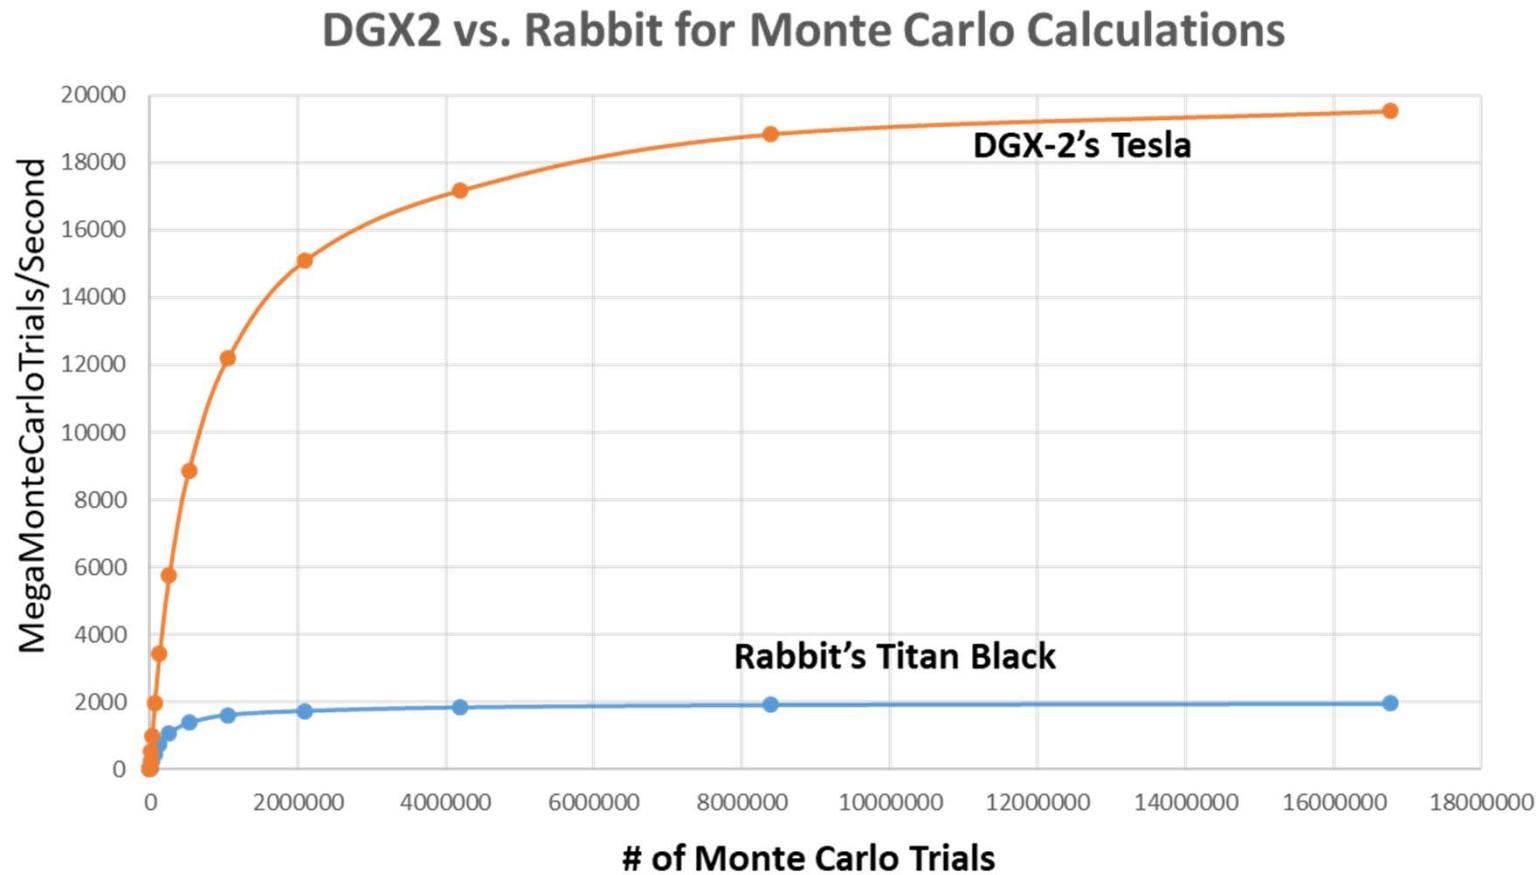
- Has 16 NVidia GPUs
- Has 28TB of disk, all SSD
- Has two 24-core Intel Xeon 8168 Platinum 2.7GHz CPUs
- Has 1.5TB of DDR4-2666 System Memory
- Runs the Rocky 9 (EL 9) Linux operating system

These are not ordinary “graphics cards”. They are Nvidia model *00 cards (V100, H100, etc.). The *00 designator mean that the GPU chips don't have any graphics capability on them, leaving more room for extra compute capabilities.

Overall compute power:

- For example, each V100 NVidia Tesla card has 5,120 CUDA Cores and 640 Tensor Cores
- This gives each 16-V100 DGX server a total of 81,920 CUDA cores and 10,240 Tensor cores
- This gives the entire 6-DGX package a total of 491,520 CUDA Cores and 61,440 Tensor Cores





BTW, you can also use our *rabbit* machine:

ssh rabbit.enr.oregonstate.edu

It is very flip-like. But, because it has a GPU in it, it is a good place to write your code and debug it. Because a lot of us will be sharing it, it is not necessarily a good place to do the final run of your code.

How to SSH to the DGX Systems

4

ssh over to a DGX submission machine --
submit-a and **submit-b** will also work

```
flip3 151% ssh submit-c.hpc.enr.oregonstate.edu
```

```
submit-c 142% module load slurm
```

Type this right away to set your path correctly



Oregon State
University
Computer Graphics

mjb – March 26, 2025

How to Check on the DGX Systems

5

submit-c 143% squeue

Check on the queues

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
3923	mime4	c_only	jayasurw	R	1-10:32:19	1	compute-e-1
3963	mime4	2Dex	jayasurw	R	16:21:03	1	compute-e-2
3876	share	CH3COOH_	chukwuk	R	1-23:36:45	1	compute-2-6
3971	nerhp	tcsh	dionnec	R	8:59:45	1	compute-h-8
3881	dgx2	bash	heli	R	1-22:50:44	1	compute-dgx2-1
3965	dgx2	bash	chenju3	R	13:47:36	1	compute-dgx2-4
3645	dgx2	bash	mishrash	R	5-16:48:09	1	compute-dgx2-5
3585	dgx2	bash	azieren	R	6-17:34:00	1	compute-dgx2-3
3583	dgx2	bash	azieren	R	6-18:26:44	1	compute-dgx2-3

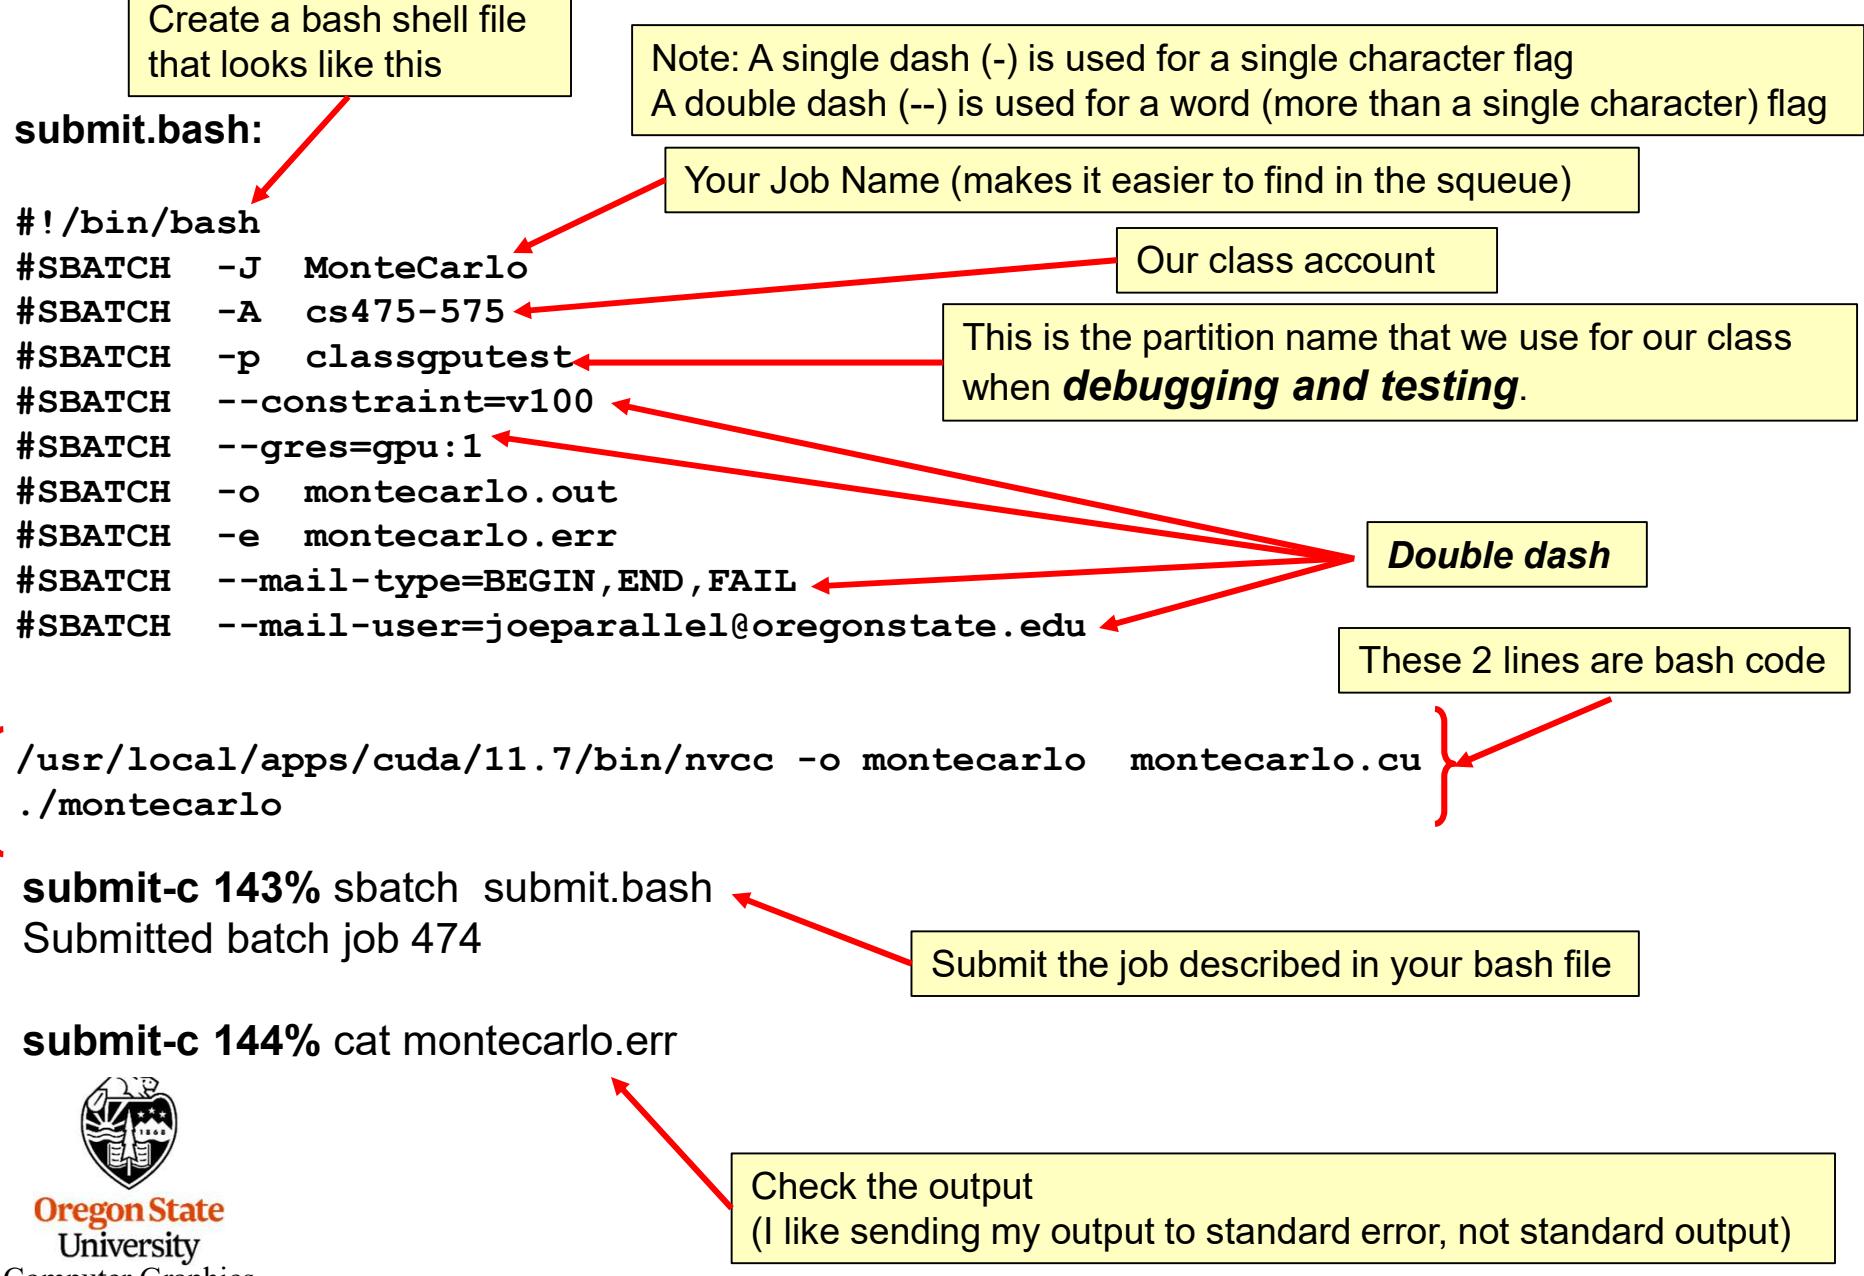
submit-c 144% sinfo

System Information

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
share*	up	7-00:00:00	2	drain	compute-4-[3-4]
share*	up	7-00:00:00	1	mix	compute-2-6
sharegpu	up	7-00:00:00	1	mix	compute-dgxs-1
sharegpu	up	7-00:00:00	3	idle	compute-dgxs-[2-3],compute-gpu
dgx2	up	7-00:00:00	1	drain	compute-dgx2-2
dgx2	up	7-00:00:00	5	mix	compute-dgx2-[1,3-6]
gpu	up	7-00:00:00	2	mix	compute-gpu[3-4]
gpu	up	7-00:00:00	1	idle	compute-gpu2
gpu	up	7-00:00:00	1	down	compute-gpu1
dgx	up	7-00:00:00	3	mix	compute-dgx2-[4-6]
dgxs	up	7-00:00:00	1	mix	compute-dgxs-1
dgxs	up	7-00:00:00	2	idle	compute-dgxs-[2-3]
class	up	1:00:00	1	mix	compute-dgxs-1
class	up	1:00:00	2	idle	compute-dgxs-[2-3]
eeecs	up	7-00:00:00	1	mix	compute-2-6

Submitting a Batch CUDA job to the DGX System

6



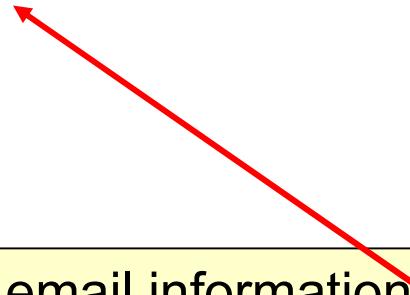
What is the Difference Between the Partitions *classgpu test* and *classgpu final*?

classgpu test lets your program get into the system sooner, but it might be running alongside other jobs, so its performance might suffer. But you don't care because you are just debugging and testing, not taking performance numbers for your report.

classgpu final makes your program wait in line until it can get dedicated resources so that you get performance results that are much more representative of what the machine can do, and thus are worthy to be listed in your report.



```
#SBATCH --mail-user=joeparallel@oregonstate.edu
```



You don't have to ask the system to email information to you, but if you do,
please be sure you spell your own email address correctly!

Our IT people are getting *really* tired of fielding the bounced emails when
people misspell their own email address.



Oregon State
University
Computer Graphics

What Showed up in my Email (which I spelled correctly)

9

From	Subject
Slurm workload manager	Slurm Job_id=3980 Name=MatrixMul Ended, Run time 00:00:12, COMPLETED, ExitCode 0 2
Slurm workload manager	Slurm Job_id=3980 Name=MatrixMul Began, Queued time 00:00:01 1



Oregon State
University
Computer Graphics

mjb – March 26, 2025

Submitting a Loop

10

submitloop.bash:

```
#!/bin/bash
#SBATCH -J MonteCarlo
#SBATCH -A cs475-575
#SBATCH -p classgpufinal
#SBATCH --constraint=v100
#SBATCH --gres=gpu:1
#SBATCH -o montecarlo.out
#SBATCH -e montecarlo.err
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mail-user=joeparallel@oregonstate.edu
```

These lines are bash code

```
{ for t in 2048 8192 131072 2097152
do
  for b in 8 16 32 64 128 256
  do
    /usr/local/apps/cuda/11.7/bin/nvcc -DNUMTRIALS=$t -DBLOCKSIZE=$b -o montecarlo montecarlo.cu
    ./montecarlo
  done
done}
```

submit-c 153% sbatch submitloop.bash
Submitted batch job 475

submit-c 154% tail -f montecarlo.err

Displays the latest output added to montecarlo.err
Keeps doing it forever.

Control-c to get out of it.



Oregon State
University
Computer Graphics

Use slurm's *scancel* if your Job Needs to Be Killed

11

submit-c 163% sbatch submitloop.bash
Submitted batch job **476**

submit-c 164% scancel 476



Oregon State
University
Computer Graphics

mjb – March 26, 2025

submit.bash:

```
#!/bin/bash
#SBATCH -J PrintInfo
#SBATCH -A cs475-575
#SBATCH -p classgpufinal
#SBATCH --constraint=v100
#SBATCH --gres=gpu:1
#SBATCH -o printinfo.out
#SBATCH -e printinfo.err
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mail-user=joeparallel@oregonstate.edu

{ module load cuda/11.7
g++ -o printinfo printinfo.cpp /usr/local/apps/cuda/11.7/lib64/libOpenCL.so.1.1 -lm -fopenmp
./printinfo }
```

Here's what *printinfo* got on one graphics card on the DGX System

13

Number of Platforms = 1

Platform #0:

Name = 'NVIDIA CUDA'

Vendor = 'NVIDIA Corporation'

Version = OpenCL 1.2 CUDA 11.2.153'

Profile = 'FULL_PROFILE'

Number of Devices = 1

Device #0:

Type = 0x0004 = CL_DEVICE_TYPE_GPU

Device Vendor ID = 0x10de (NVIDIA)

Device Maximum **Compute Units** = 80

Device Maximum Work Item Dimensions = 3

Device Maximum Work Item Sizes = 1024 x 1024 x 64

Device Maximum Work Group Size = 1024

Device Maximum **Clock Frequency** = 1530 MHz

For comparison, *rabbit*'s graphics card has **15** Compute Units

Device Extensions:

cl_khr_global_int32_base_atomics

cl_khr_global_int32_extended_atomics

cl_khr_local_int32_base_atomics

cl_khr_local_int32_extended_atomics

cl_khr_fp64

cl_khr_byte_addressable_store

cl_khr_icd

cl_khr_gl_sharing

cl_nv_compiler_options

cl_nv_device_attribute_query

cl_nv_pragma_unroll

cl_nv_copy_opts

cl_nv_create_buffer

