

The Message Passing Interface (MPI): Parallelism on Distributed CPUs

<http://mpi-forum.org>
<https://www.open-mpi.org/>



Oregon State
University
Mike Bailey

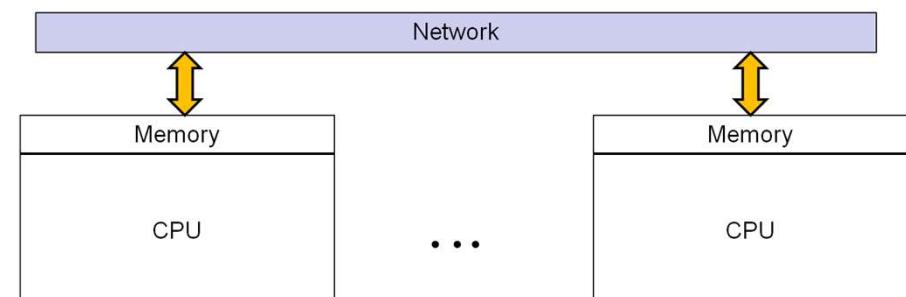
mjb@cs.oregonstate.edu



This work is licensed under a [Creative Commons](#)
[Attribution-NonCommercial-NoDerivatives 4.0](#)
[International License](#)



Oregon State
University
Computer Graphics



Why Two URLs?



<http://mpi-forum.org>

This is the definitive reference for the MPI standard. Go here if you want to read the official specification, which, BTW, continues to evolve.

<https://www.open-mpi.org/>

This consortium formed later. This is the open source version of MPI. If you want to start using MPI, I recommend you look here. This is the MPI that the COE systems use

<https://www.open-mpi.org/doc/v4.0/>

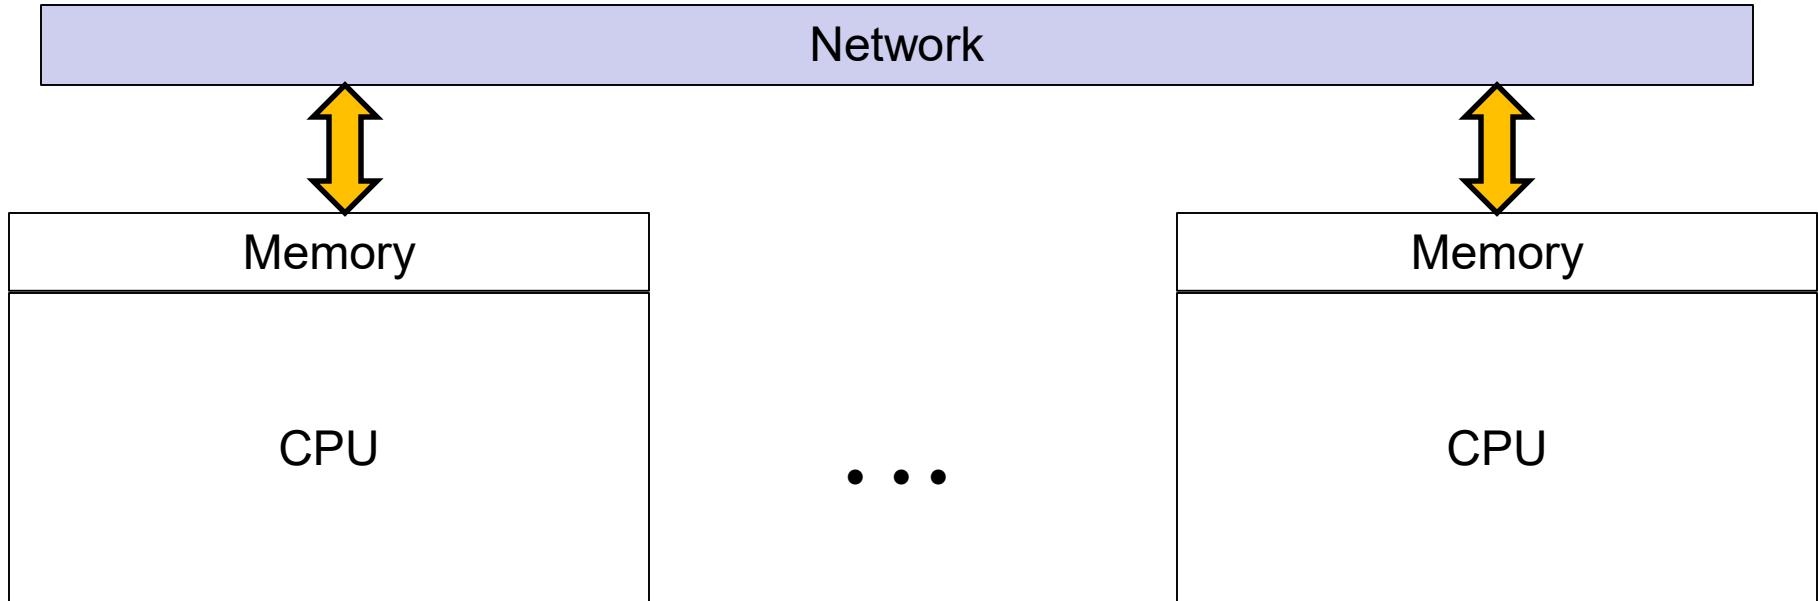
This URL is also really good – it is a link to all of the MPI man pages



The Open MPI Consortium



MPI: The Basic Idea



Programs on different CPUs coordinate computations by passing messages between each other

Note: Each CPU in the MPI “cluster” must be prepared ahead of time by having the MPI server code installed on it. It must then have that server code running and listening on its socket connection.

Each MPI CPU must also have an integer ID assigned to it (called its **rank**).

This is how modern supercomputers work!

5



Oregon State
University
Computer Graphics

The Texas Advanced Computing Center's *Frontera* supercomputer

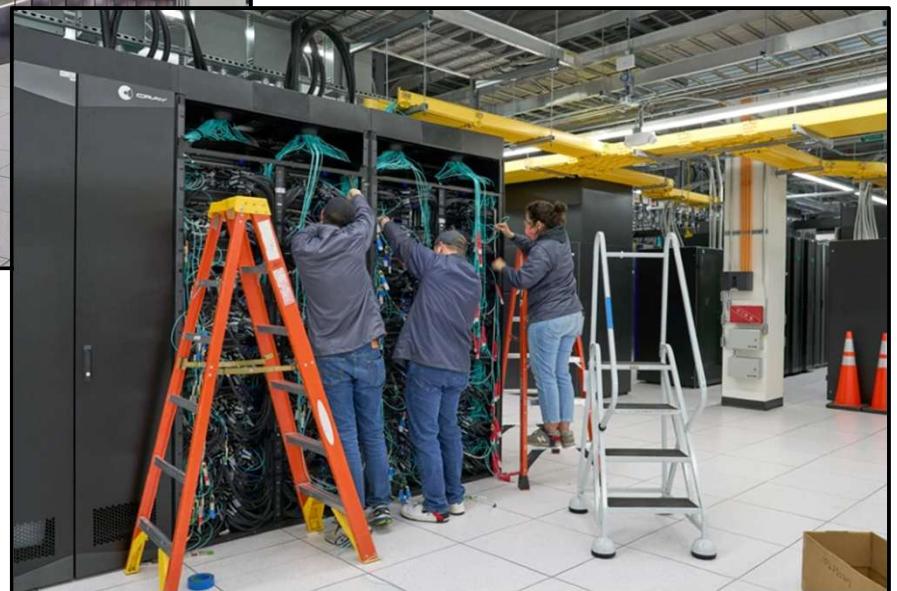
mjb – March 26, 2025

This is how modern supercomputers work!

6



- 50,000 processors
- 38,000 GPUs
- >1 exaflop (10^{18})
- 27 MW



Oregon State
University
Computer Graphics

The Oakridge National Lab *Frontier* supercomputer

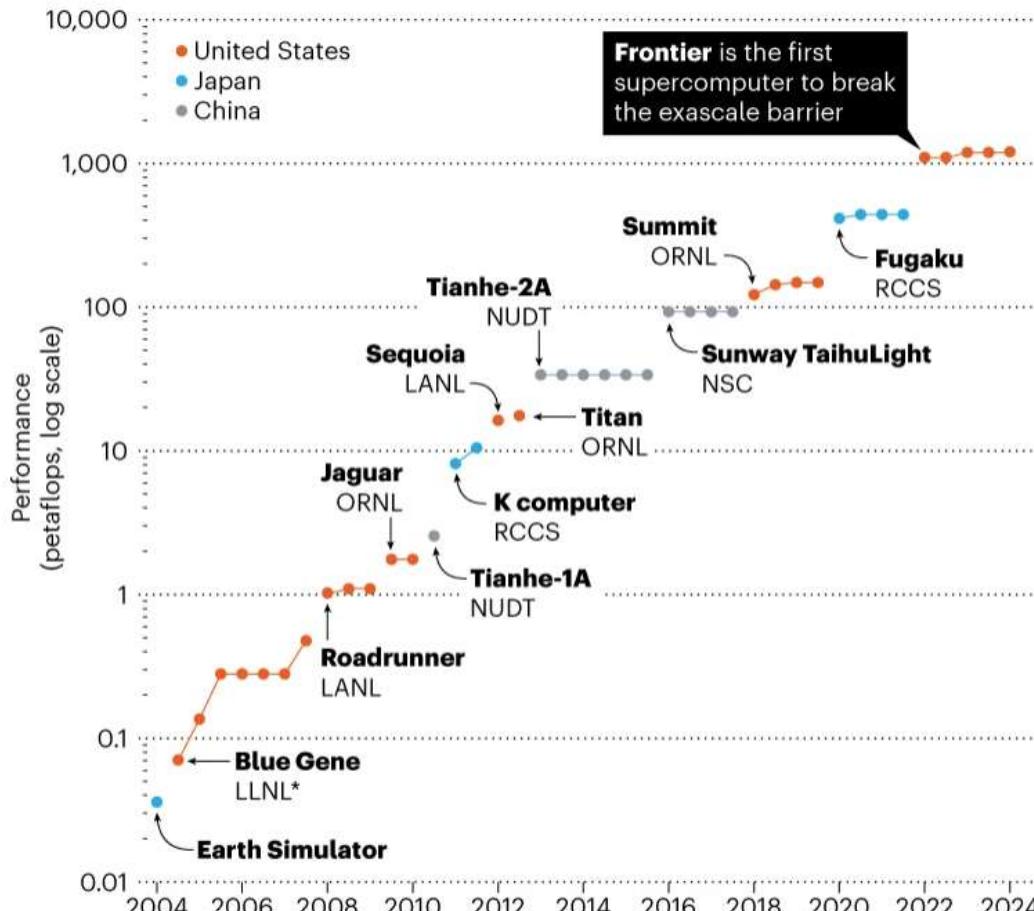
mjb – March 26, 2025

This paradigm is how modern supercomputers work!

7

SPEED RECORDS

Design advances over the past two decades have produced huge increases in speed for the world's fastest supercomputers, which are rated using a benchmark test to generate the TOP500 list.



*LLNL, Lawrence Livermore National Laboratory; LANL, Los Alamos National Lab; NSC, National Supercomputer Center; NUDT, National University of Defense Technology; ORNL, Oak Ridge National Lab; RCCS, Riken Center for Computational Science (previously Advanced Institute for Computational Science).

How to SSH to the COE MPI Cluster

8

ssh over to an MPI submission machine --
submit-a and **submit-b** will also work

```
flip3 151% ssh submit-c.hpc.engr.oregonstate.edu
```

```
submit-c 142% module load slurm  
submit-c 143% module load openmpi
```

Type these right away to set your path correctly

BTW, you can find out more about the COE cluster here:

<https://it.engineering.oregonstate.edu/hpc>

The College of Engineering HPC cluster is a heterogeneous mix of over 130 servers providing 4,800 CPU cores, over 230 GPUs, and over 50 TB of total RAM. The systems are connected via gigabit ethernet and Infiniband. Most of the latest servers utilize Mellanox EDR or HDR InfiniBand network connection. The cluster also has access to 500TB global scratch space. The CoE HPC Cluster is rated at over 2,200 peak TFLOPS (double-precision).

Compiling and Running

```
mpicc -o program program.c . . .
```

← C

or

```
mpic++ -o program program.cpp . . .
```

← C++

```
mpiexec -mca btl self,tcp -np 4 program
```

All distributed CPUs execute the
same program at the same time

of CPUs (processors) to use

Warning – use *mpic++* and *mpiexec*!

Don't use g++.

Don't run by just typing the name of the executable!

Running with a *bash* Batch Script

submit.bash:

```
#!/bin/bash
#SBATCH -J Fourier
#SBATCH -A cs475-575
#SBATCH -p classmpitest
#SBATCH -N 8    # number of nodes
#SBATCH -n 8    # number of tasks
#SBATCH -o mpiproject.out
#SBATCH -e mpiproject.err
#SBATCH --mail-type=END,FAIL
#SBATCH --mail-user=joeparallel@oregonstate.edu
module load openmpi
mpic++ mpiproject.cpp -o mpiproject -lm
mpiexec -mca btl self,tcp -np 4 ./mpiproject
```

Your Job Name

This is the partition name that we use for our class when debugging and testing your program. Use **classmpifinal** for taking your final performance numbers.

These 3 lines are bash code



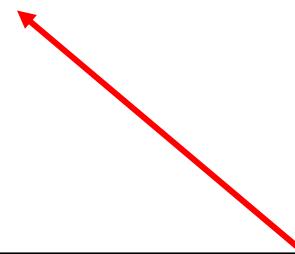
What is the Difference Between the Partitions *classmpitest* and *classmpifinal*?

classmpitest lets your program get into the system sooner, but it might be running alongside other jobs, so its performance might suffer. But, you don't care because you are just compiling and debugging, not taking performance numbers for your report.

classmpifinal makes your program wait in line until it can get dedicated resources so that you get performance results that are much more representative of what the machines can do, and thus are worthy to be listed in your report.



```
#SBATCH --mail-user=joeparallel@oregonstate.edu
```



You don't have to ask the system to email information to you, but if you do,
please be sure you spell your own email address correctly!

Our IT people are getting *really* tired of fielding the bounced emails when
people misspell their own email address.



Oregon State
University
Computer Graphics

Use slurm's *scancel* if your Job Needs to Be Killed

13

submit-c 143% sbatch submit.bash
Submitted batch job 258759

submit-c 144% scancel 258759

Setting Up and Finishing MPI

```
#include <mpi.h>

int
main( int argc, char *argv[ ] )
{
    . . .
    MPI_Init( &argc, &argv );
    . . .
    MPI_Finalize( );
    return 0;
}
```

You don't need to process command line arguments if you don't want to.
You can just call MPI_Init() as:

MPI_Init(NULL, NULL);

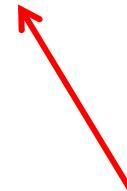
MPI Follows a Single-Program-Multiple-Data (SPMD) Model

A **communicator** is a collection of CPUs that are capable of sending messages to each other

Oh, look, a
communicator of
Corvallis deer!



Oh, look, a
communicator of
Corvallis turkeys!



This requires MPI server code getting installed on all those CPUs. That code then needs to be running and listening on a socket connection. Only an administrator can do this.

Getting information about our place in the **communicator**:

```
int numCPUs;      // total # of cpus involved
int me;           // which one I am

MPI_Comm_size( MPI_COMM_WORLD, &numCPUs );

MPI_Comm_rank( MPI_COMM_WORLD, &me );
```



Size, i.e., how many altogether?

Rank, i.e., which one am I?

It is then each CPU's job to figure out what piece of the overall problem it is responsible for and then go do it.

A First Test of MPI

```
#include <stdio.h>
#include <math.h>
#include <mpi.h>

#define THEBOSS 0

int
main( int argc, char *argv[ ] )
{
    MPI_Init( &argc, &argv );

    int numCPUs;      // total # of cpus involved
    int me;           // which one I am

    MPI_Comm_size( MPI_COMM_WORLD, &numCPUs );
    MPI_Comm_rank( MPI_COMM_WORLD, &me );

    if( me == THEBOSS )
        fprintf( stderr, "Rank %d says that we have a Communicator of size %d\n", THEBOSS, numCPUs );
    else
        fprintf( stderr, "Welcome from Rank %d\n", me );

    MPI_Finalize( );
    return 0;
}
```

submit-c 165% mpiexec -np 16 ./first

Welcome from Rank 13
Welcome from Rank 15
Welcome from Rank 3
Welcome from Rank 7
Welcome from Rank 5
Welcome from Rank 8
Welcome from Rank 9
Welcome from Rank 11
Rank 0 says that we have a Communicator of size 16
Welcome from Rank 1
Welcome from Rank 12
Welcome from Rank 14
Welcome from Rank 6
Welcome from Rank 2
Welcome from Rank 10
Welcome from Rank 4

submit-c 166% mpiexec -np 16 ./first

Welcome from Rank 1
Welcome from Rank 5
Welcome from Rank 7
Welcome from Rank 9
Welcome from Rank 11
Welcome from Rank 13
Welcome from Rank 15
Rank 0 says that we have a Communicator of size 16
Welcome from Rank 2
Welcome from Rank 3
Welcome from Rank 4
Welcome from Rank 6
Welcome from Rank 8
Welcome from Rank 12
Welcome from Rank 14
Welcome from Rank 10

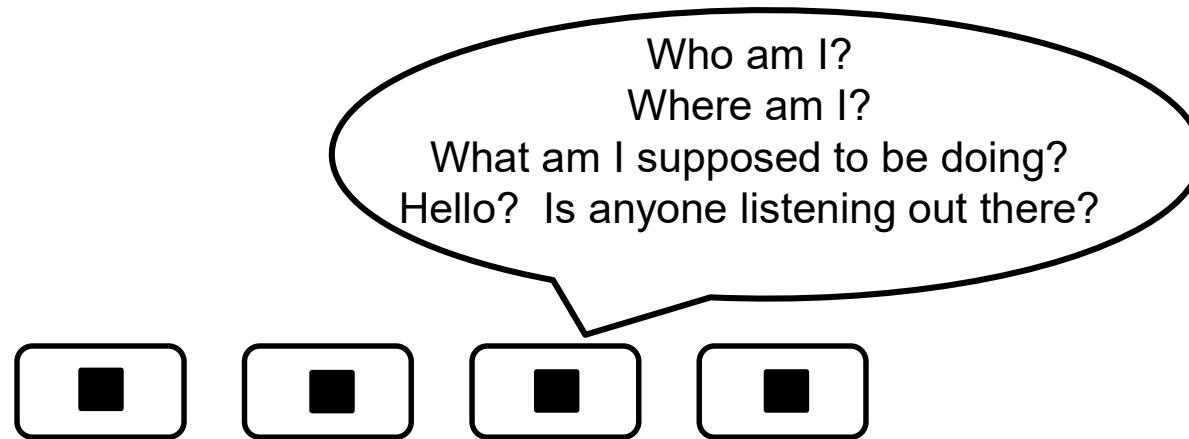
submit-c 167% mpiexec -np 16 ./first

Welcome from Rank 9
Welcome from Rank 11
Welcome from Rank 13
Welcome from Rank 7
Welcome from Rank 1
Welcome from Rank 3
Welcome from Rank 10
Welcome from Rank 15
Welcome from Rank 4
Welcome from Rank 5
Rank 0 says that we have a Communicator of size 16
Welcome from Rank 2
Welcome from Rank 6
Welcome from Rank 8
Welcome from Rank 14
Welcome from Rank 12

submit-c 168% mpiexec -np 16 ./first

Welcome from Rank 13
Welcome from Rank 15
Welcome from Rank 7
Welcome from Rank 3
Welcome from Rank 5
Welcome from Rank 9
Welcome from Rank 11
Welcome from Rank 1
Welcome from Rank 12
Welcome from Rank 14
Welcome from Rank 4
Welcome from Rank 2
Rank 0 says that we have a Communicator of size 16
Welcome from Rank 8
Welcome from Rank 10
Welcome from Rank 6

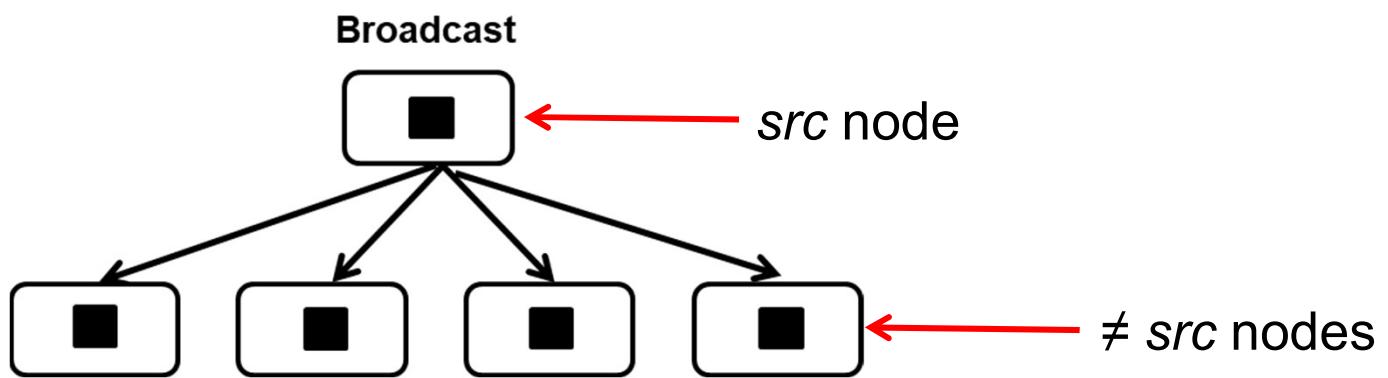
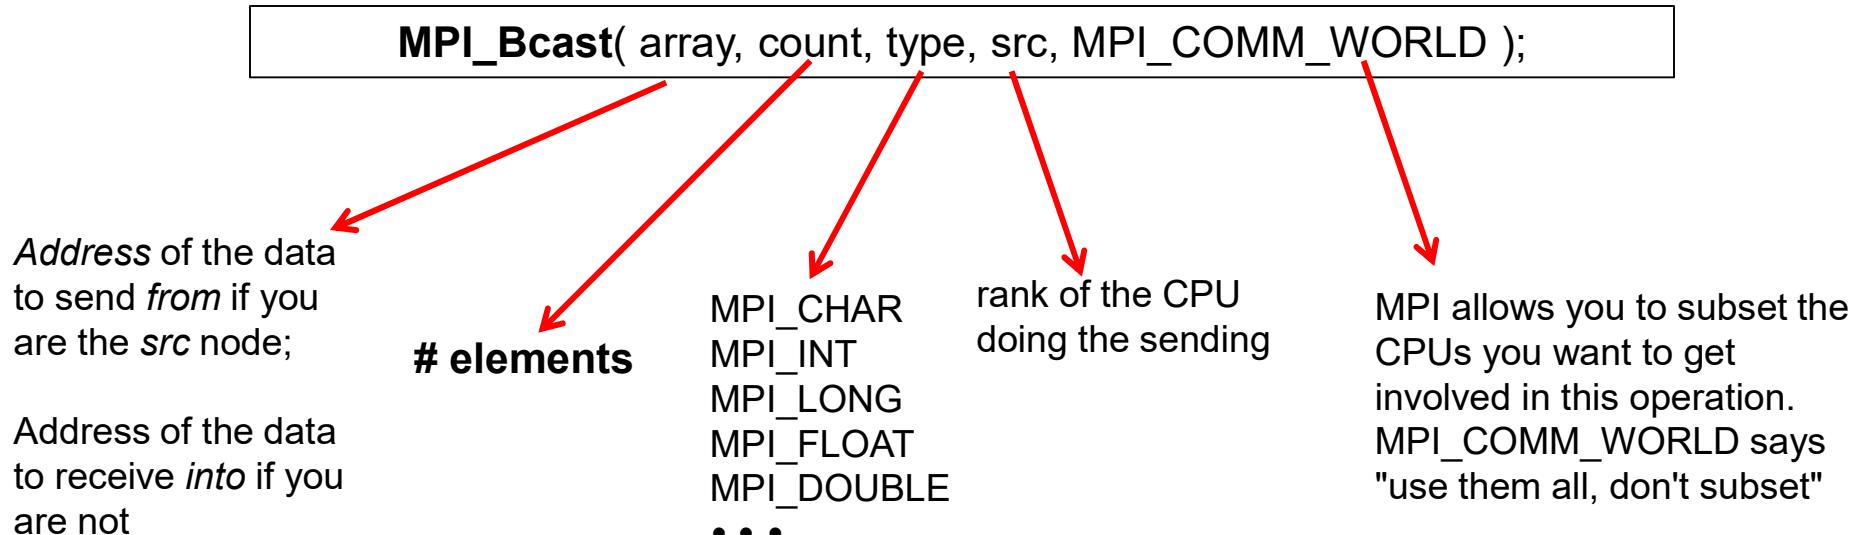
So, we have a group (a “communicator”) of distributed CPUs.
How do they communicate about what work they are supposed to do?



Example: You could coordinate the units of our DGX system using MPI



A Good Place to Start: MPI Broadcasting



MPI Broadcast Example

This is our heat transfer equation from before. Clearly, every CPU will need to know this value.

$$\Delta T_i = \left(\frac{k}{\rho C} \right) \left(\frac{T_{i-1} - 2T_i + T_{i+1}}{(\Delta x)^2} \right) \Delta t$$

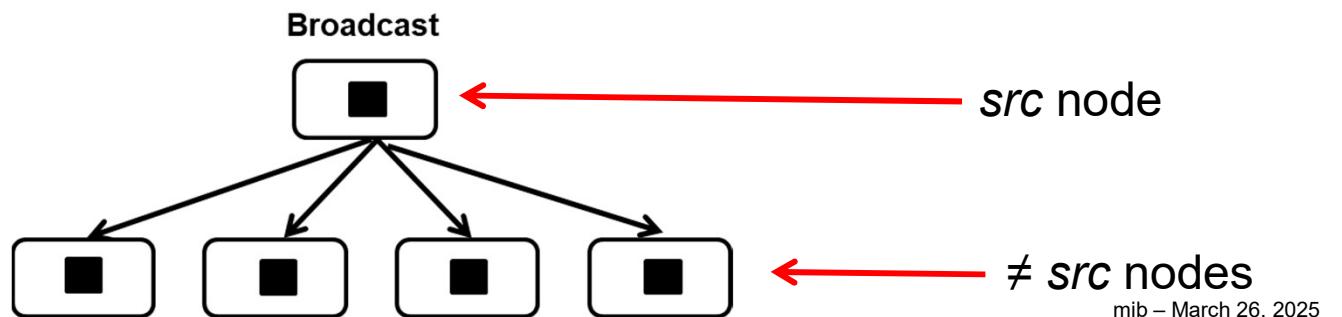
```

int numCPUs;
int me;
float k_over_rho_c; // the THEBOSS node will know this value, the others won't (yet)

#define THEBOSS 0

MPI_Comm_size( MPI_COMM_WORLD, &numCPUs ); // how many are in this communicator
MPI_Comm_rank( MPI_COMM_WORLD, &me ); // which one am I?
...
if( me == THEBOSS ) // I am the THEBOSS: this identifies this call as a send
{
    << read k_over_rho_c from the data file >>
}
MPI_Bcast( &k_over_rho_c, 1, MPI_FLOAT, THEBOSS, MPI_COMM_WORLD ); // send if I am THEBOSS, and receive if not

```



How Does this Work?
Think Star Trek Wormholes!



Sending Data from One Source CPU to Just One Destination CPU

```
MPI_Send( array, numToSend, type, dst, tag, MPI_COMM_WORLD );
```

address of data to send from

elements

(note: this is the number of *elements*, not the number of *bytes*!)

MPI_CHAR
MPI_INT
MPI_LONG
MPI_FLOAT
MPI_DOUBLE
• • •

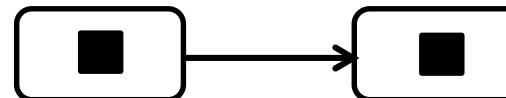
rank of the CPU to send to

An integer or character to differentiate this transmission from any other transmission.
I like to use chars.

MPI allows you to subset the CPUs you want to get involved in this operation.
MPI_COMM_WORLD says "use them all, don't subset"

Rules:

- One message from a specific *src* to a specific *dst* cannot overtake a previous message from the same *src* to the same *dst*.
- MPI_Send() blocks until the transfer is far enough along that *array* can be destroyed or re-used.
- There are no guarantees on order from different *src*'s .



```
MPI_Recv( array, maxCanReceive, type, src, tag, MPI_COMM_WORLD, &status );
```

address of data to receive into

elements we can receive, at most

MPI_CHAR
MPI_INT
MPI_LONG
MPI_FLOAT
MPI_DOUBLE
• • •

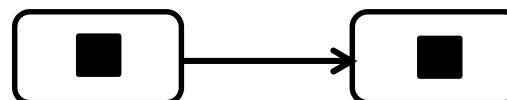
Rank of the CPU we are expecting to get a transmission from

Type = MPI_Status

An integer or character to differentiate what transmission we are looking for with this call (be sure this matches what the sender is sending!). I like to use chars.

Rules:

- The receiver blocks waiting for data that matches what it declares to be looking for
- One message from a specific *src* to a specific *dst* cannot overtake a previous message from the same *src* to the same *dst*
- There are no guarantees on the order from different *src*'s
- The order from different *src*'s could be implied in the *tag*
- **status** is type MPI_Status – the “&status” can be replaced with MPI_STATUS_IGNORE



Example

Remember, this *identical code* runs on all CPUs:

```

int numCPUs;
int me;
#define MYDATA_SIZE 128
char mydata[ MYDATA_SIZE ];
#define THEBOSS 0

MPI_Comm_size( MPI_COMM_WORLD, &numCPUs );
MPI_Comm_rank( MPI_COMM_WORLD, &me );

if( me == THEBOSS ) // the primary
{
    for( int dst = 0; dst < numCPUs; dst++ )
    {
        if( dst != THEBOSS )
        {
            char *InputData = "Hello, Beavers!";
            MPI_Send( InputData, strlen(InputData)+1, MPI_CHAR, dst, 'B', MPI_COMM_WORLD );
        }
    }
}
else // a secondary
{
    MPI_Recv( myData, MYDATA_SIZE, MPI_CHAR, THEBOSS, 'B', MPI_COMM_WORLD, MPI_STATUS_IGNORE );
    printf( "%s from rank # %d\n", in, me );
}

```

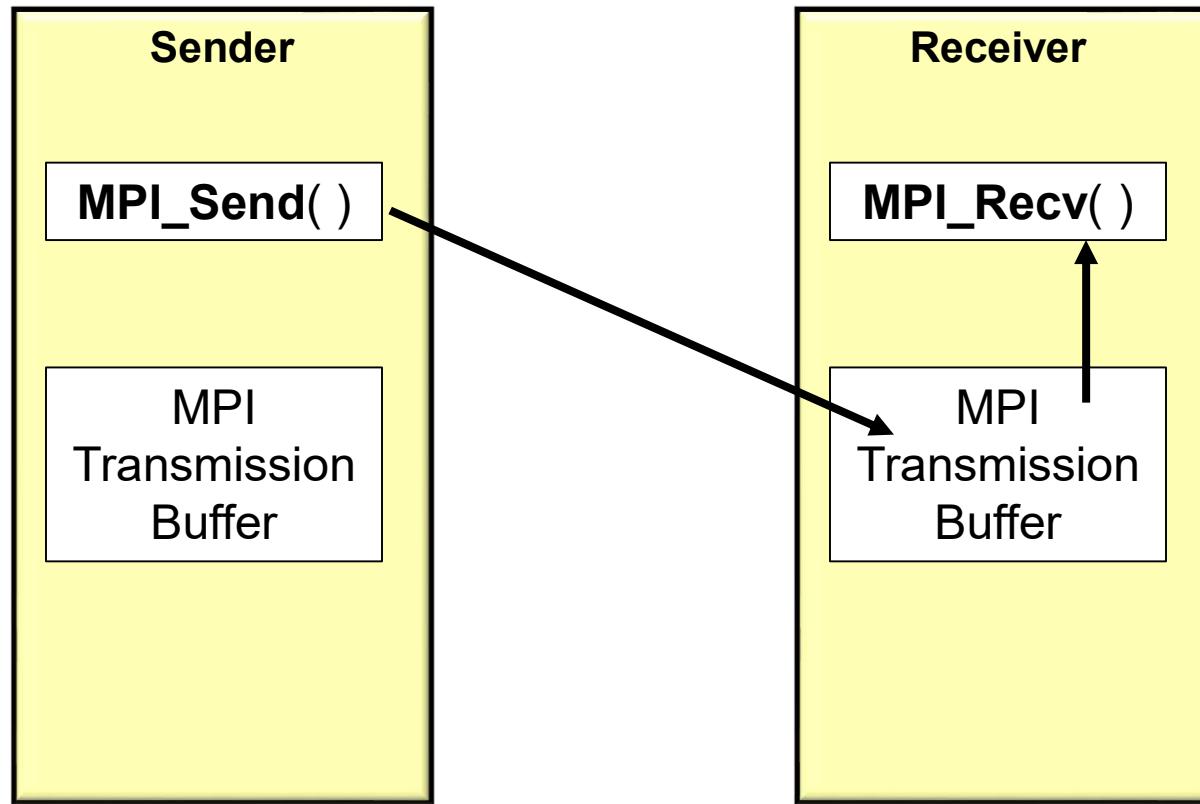
Be sure the receiving tag matches the sending tag

The tag to expect

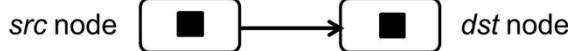
The tag to label this transmission with

Or You are highly discouraged from sending to yourself. Because both the send and receive are capable of blocking, the result could be deadlock.

How does MPI let the Sender perform an MPI_Send() even if the Receivers are not ready to MPI_Recv()?



MPI_Send() blocks until the transfer is far enough along that the *array* can be destroyed or re-used.



Another Example

You typically don't send the entire workload to each dst – you just send part of it, like this:

```
#define NUMELEMENTS  ??????
int numCPUs;
int me;
#define THEBOSS    0

MPI_Comm_size( MPI_COMM_WORLD, &numCPUs );
MPI_Comm_rank( MPI_COMM_WORLD, &me );

int PPSIZE = NUMELEMENTS / numCPUs;      // per-processor data size -- assuming it comes out evenly
float *myData = new float [ PPSIZE ];

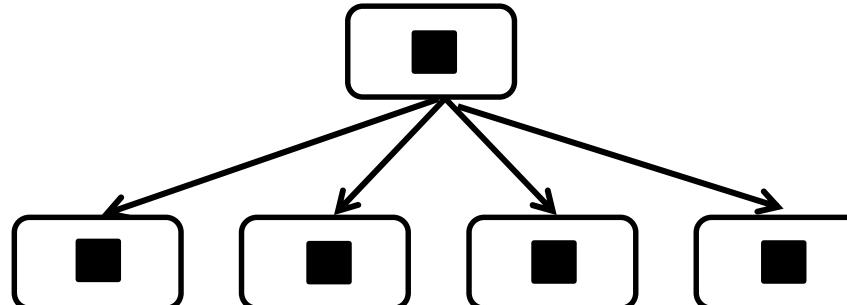
if( me == THEBOSS ) // the sender
{
    float *InputData = new float [ NUMELEMENTS ];
    << read the full input data into InputData from disk >>
    for( int dst = 0; dst < numCPUs; dst++ )
    {
        if( dst != THEBOSS )
        {
            MPI_Send( &InputData[dst*PPSIZE], PPSIZE, MPI_FLOAT, dst, 0, MPI_COMM_WORLD );
        }
    }
}
else // a receiver
{
    MPI_Recv( myData, PPSIZE, MPI_FLOAT, THEBOSS, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE );
    // do something with this subset of the data
}
```

The address of node *dst*'s share of the data to send

Each *dst* node will store its data in this array

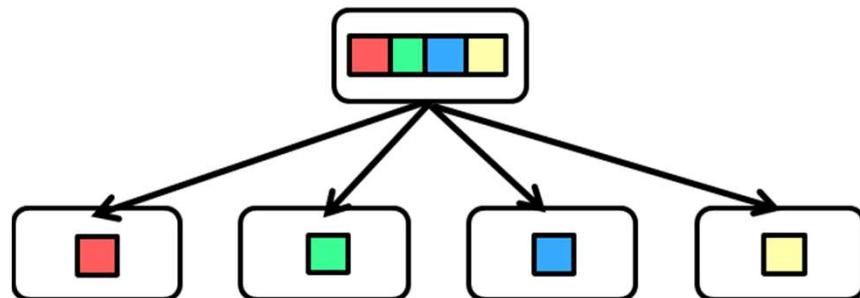
In Distributed Computing, You Often Hear About These Design Patterns²⁷

Broadcast



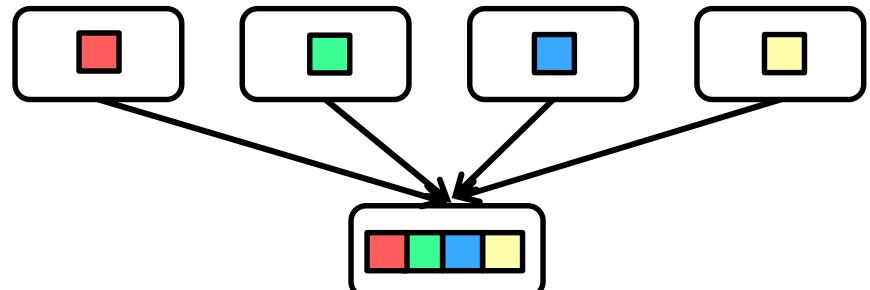
This sends the same data to each CPU

Scatter



This sends pieces of the data to each CPU

Gather

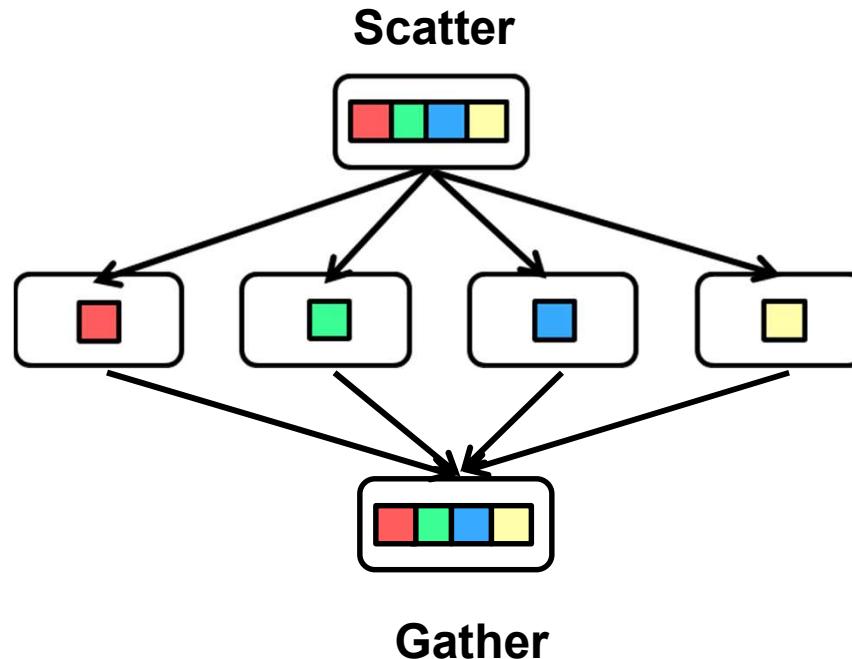


This brings back pieces of the results from each CPU



Scatter and Gather Usually Go Together

This sends pieces of the data to each CPU



Note surprisingly, this is referred to by the combined term **Scatter/Gather**

MPI Scatter

Take a data array, break it into ~equal portions, and send it to each CPU

```
MPI_Scatter( snd_array, snd_count, snd_type, rcv_array, rcv_count, rcv_type, src, MPI_COMM_WORLD );
```

The total large array
to split up

elements to send
per-processor

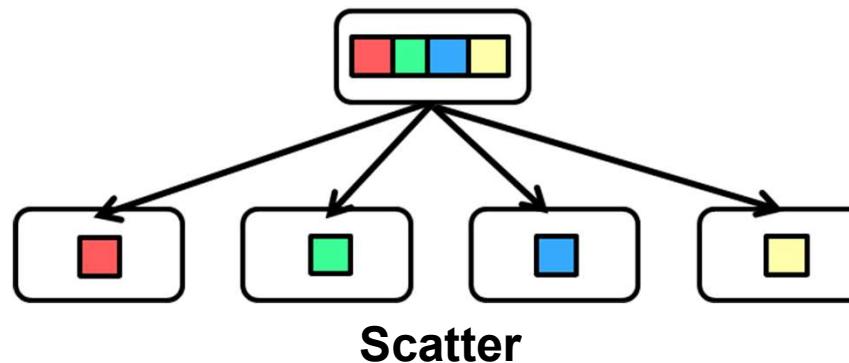
MPI_CHAR
MPI_INT
MPI_LONG
MPI_FLOAT
MPI_DOUBLE
• • •

Local array to store
this CPU's piece in

elements to receive
per-processor

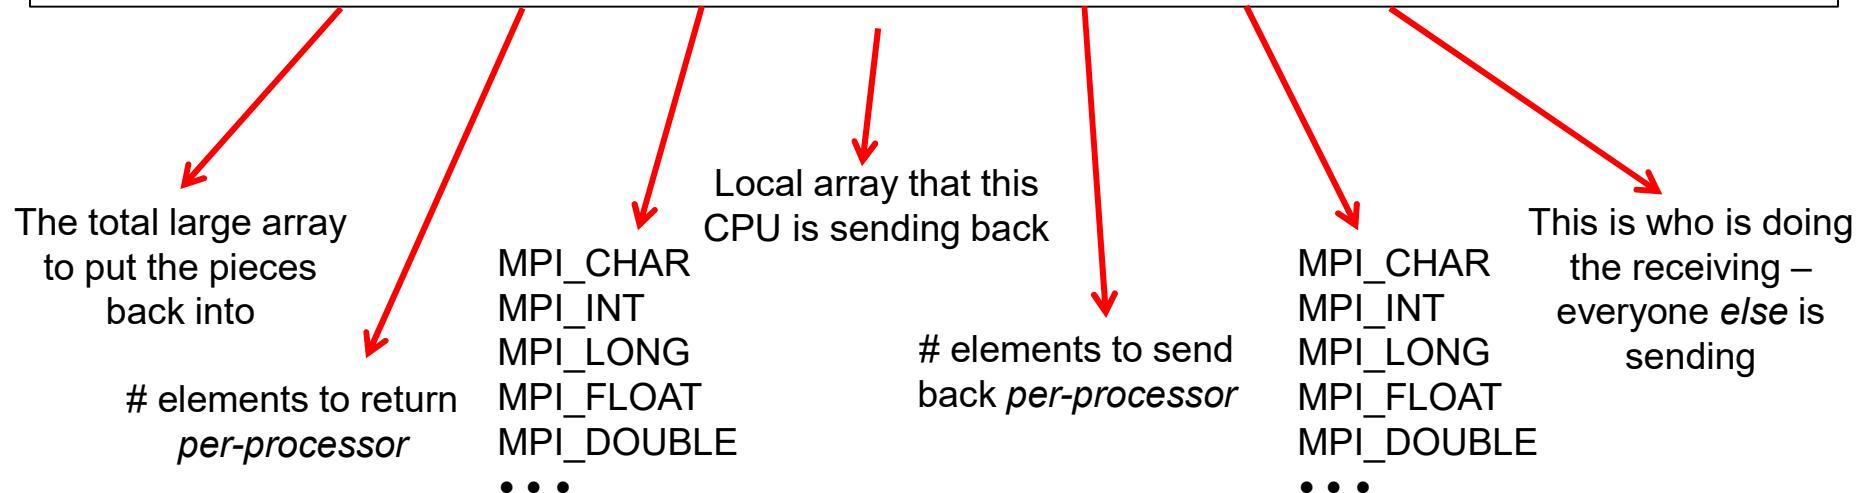
MPI_CHAR
MPI_INT
MPI_LONG
MPI_FLOAT
MPI_DOUBLE
• • •

This is who is doing
the sending –
everyone else is
receiving

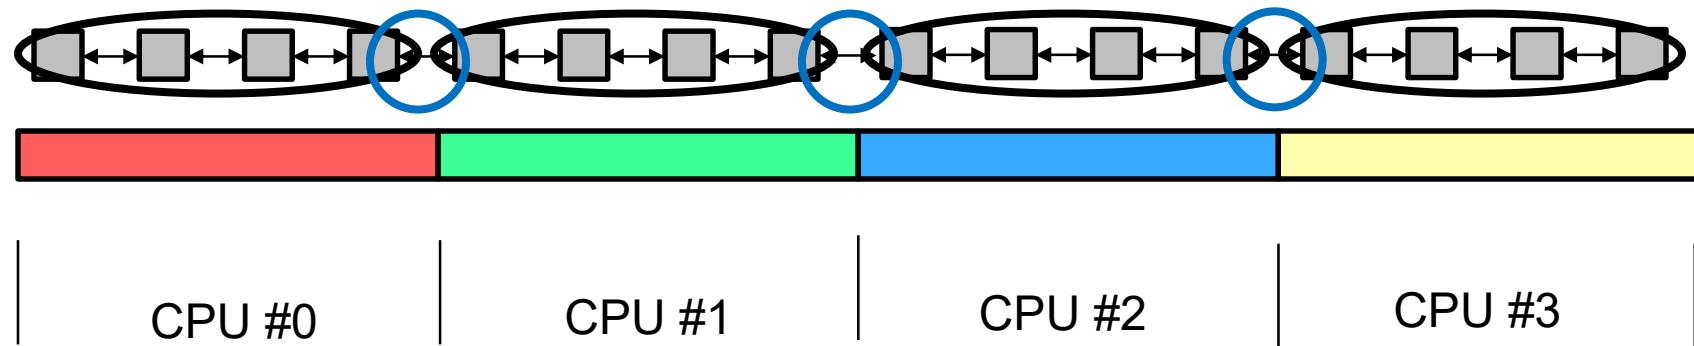


MPI Gather

```
MPI_Gather( snd_array, snd_count, snd_type, rcv_array, rcv_count, rcv_type, dst, MPI_COMM_WORLD );
```



Remember This? It's Baaaaack as a complete Scatter/Gather Example



The **Compute : Communicate Ratio** still applies, except that it is even more important now because there is much more overhead in the Communicate portion.

This pattern of breaking a big problem up into pieces, sending them to different CPUs, computing on the pieces, and getting the results back is *very* common. That's why MPI has its own scatter and gather functions.



heat.cpp, I

```
#include <stdio.h>
#include <math.h>
#include <mpi.h>

const float RHO = 8050.;
const float C  = 0.466;
const float K  = 20.;

float k_over_rho_c = K / (RHO*C);           // units of m^2/sec NOTE: this cannot be a const!
// K / (RHO*C) = 5.33x10^-6 m^2/sec

const float DX  =      1.0;
const float DT  =      1.0;

#define THEBOSS 0

#define NUMELEMENTS      (8*1024*1024)
#define NUM_TIME_STEPS   4
#define DEBUG            false

float *    NextTemps;           // per-processor array to hold computer next-values
int       NumCpus;             // total # of cpus involved
int       PPSIZE;              // per-processor local array size
float *   PPTemps;             // per-processor local array temperature data
float *   TempData;            // the overall NUMELEMENTS-big temperature data

void     DoOneTimeStep( int );
```

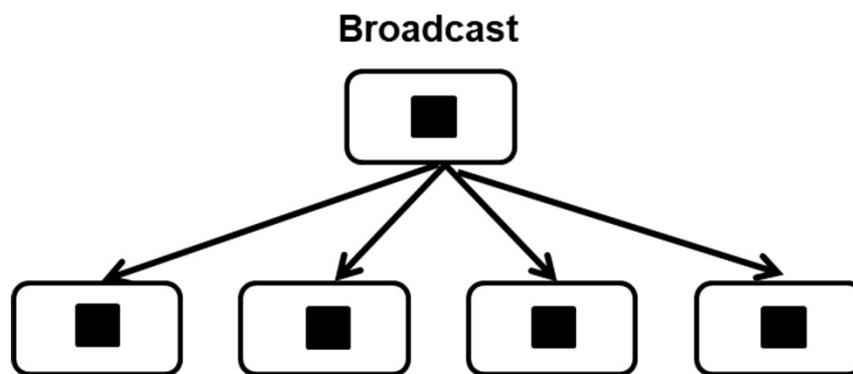
```
int
main( int argc, char *argv[ ] )
{
    MPI_Init( &argc, &argv );

    int me;           // which one I am

    MPI_Comm_size( MPI_COMM_WORLD, &NumCpus );
    MPI_Comm_rank( MPI_COMM_WORLD, &me );

    // decide how much data to send to each CPU:
    PPSIZE = NUMELEMENTS / NumCpus;           // assuming it comes out evenly
    PPTemps = new float [PPSIZE];   // all CPUs now have this uninitialized Local array
    NextTemps = new float [PPSIZE]; // all CPUs now have this uninitialized local array too

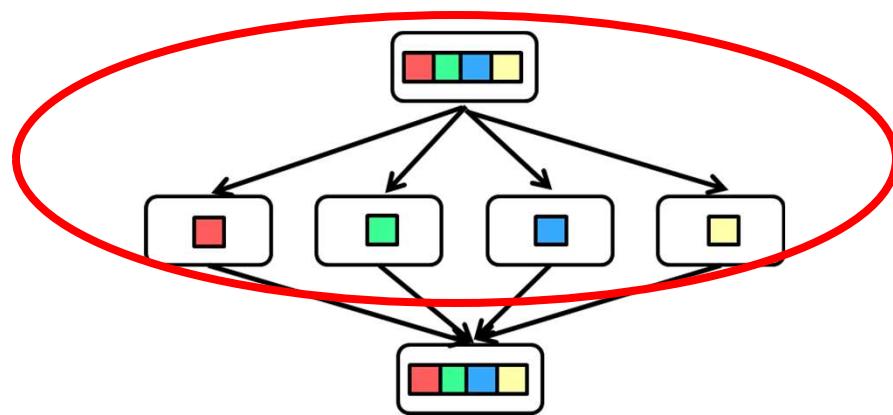
    // broadcast the constant:
    MPI_Bcast( (void *)&k_over_rho_c, 1, MPI_FLOAT, THEBOSS, MPI_COMM_WORLD );
```



heat.cpp, III

```
if( me == THEBOSS ) // this is the data-creator
{
    TempData = new float [NUMELEMENTS];
    for( int i = 0; i < NUMELEMENTS; i++ )
        TempData[ i ] = 0.;
    TempData[ NUMELEMENTS/2 ] = 100.;

MPI_Scatter( TempData, PPSIZE, MPI_FLOAT, PPTemps, PPSIZE, MPI_FLOAT,
THEBOSS, MPI_COMM_WORLD );
```



heat.cpp, IV

```

// all the PPTemps arrays have now been filled
// do the time steps:

double time0 = MPI_Wtime( );

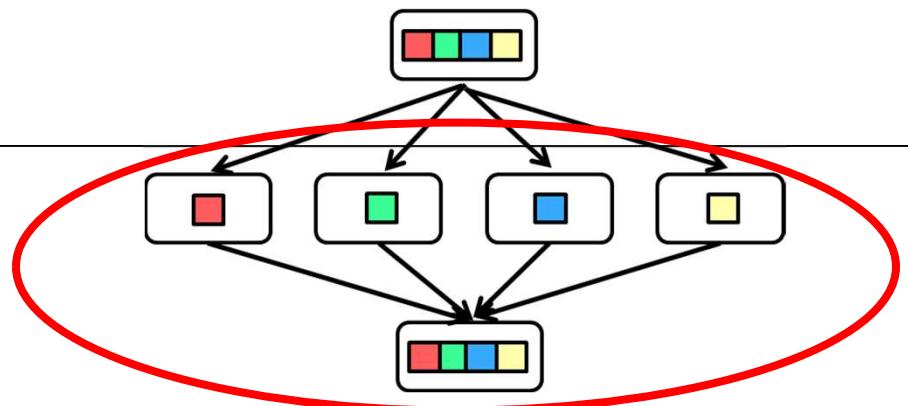
for( int steps = 0; steps < NUM_TIME_STEPS; steps++ )
{
    // do the computation for one time step:
    DoOneTimeStep( me );

    // ask for all the data:
#endif WANT_EACH_TIME_STEPS_DATA_BACK
    MPI_Gather( PPTemps, PPSIZE, MPI_FLOAT, TempData, PPSIZE, MPI_FLOAT,
                THEBOSS, MPI_COMM_WORLD );
#endif
}

#ifndef WANT_EACH_TIME_STEPS_DATA_BACK
    MPI_Gather( PPTemps, PPSIZE, MPI_FLOAT, TempData, PPSIZE, MPI_FLOAT,
                THEBOSS, MPI_COMM_WORLD );
#endif

double time1 = MPI_Wtime( );

```



heat.cpp, V

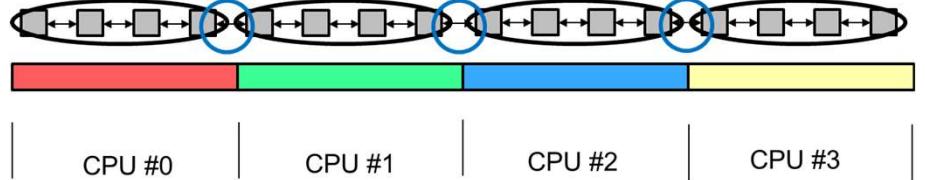
```
if( me == THEBOSS )
{
    double seconds = time1 - time0;
    double performance =
        (double)NUM_TIME_STEPS * (double)NUMELEMENTS / seconds / 1000000.;
        // mega-elements computed per second
    fprintf( stderr, "%3d, %10d, %8.2f\n", NumCpus, NUMELEMENTS, performance );
}

MPI_Finalize( );
return 0;
}
```

```
// read from PerProcessorData[ ], write into NextTemps[ ]
void
DoOneTimeStep( int me )
{
    MPI_Status status;

    // send out the left and right end values:
    // (the tag is from the point of view of the sender)
    if( me != 0 )                                // i.e., if i'm not the first group on the left
    {
        // send my PPTemps[0] to me-1 using tag 'L'
        MPI_Send( &PPTemps[0], 1, MPI_FLOAT, me-1, 'L', MPI_COMM_WORLD );
        if( DEBUG ) fprintf( stderr, "%3d sent 'L' to %3d\n", me, me-1 );
    }

    if( me != NumCpus-1 )                         // i.e., not the last group on the right
    {
        // send my PPTemps[PPSize-1] to me+1 using tag 'R'
        MPI_Send( &PPTemps[PPSize-1], 1, MPI_FLOAT, me+1, 'R', MPI_COMM_WORLD );
        if( DEBUG ) fprintf( stderr, "%3d sent 'R' to %3d\n", me, me+1 );
    }
}
```



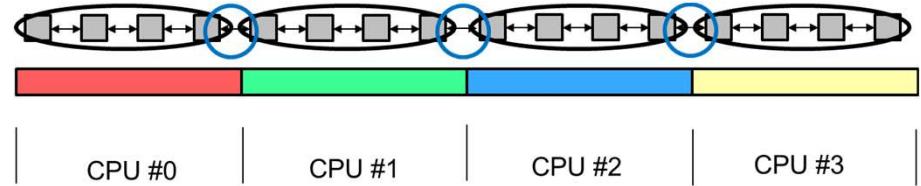
```

float left = 0.;
float right = 0.;

if( me != 0 )                                // i.e., if i'm not the first group on the left
{
    // receive my "left" from me-1 using tag 'R'
    MPI_Recv( &left, 1, MPI_FLOAT, me-1, 'R', MPI_COMM_WORLD, &status );
    if( DEBUG )  fprintf( stderr, "%3d received 'R' from %3d\n", me, me-1 );
}

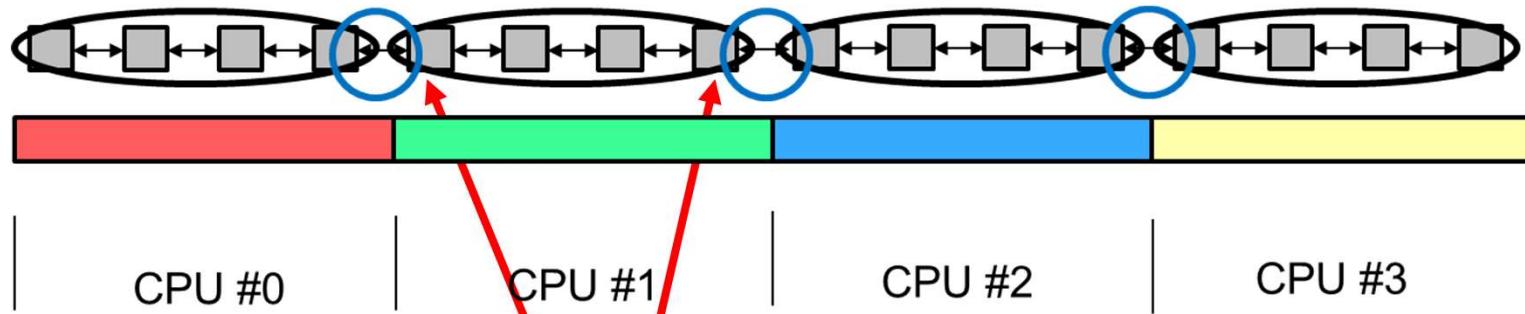
if( me != NumCpus-1 )                          // i.e., not the last group on the right
{
    // receive my "right" from me+1 using tag 'L'
    MPI_Recv( &right, 1, MPI_FLOAT, me+1, 'L', MPI_COMM_WORLD, &status );
    if( DEBUG )  fprintf( stderr, "%3d received 'L' from %3d\n", me, me+1 );
}

```



Sharing Values Across the Boundaries

39



1 sent 'L' to 0
1 sent 'R' to 2
2 sent 'L' to 1
2 sent 'R' to 3
2 received 'R' from 1
0 sent 'R' to 1
0 received 'L' from 1
1 received 'R' from 0
1 received 'L' from 2
3 sent 'L' to 2
3 received 'R' from 2
2 received 'L' from 3



DoOneTimeStep, III

```
// first element on the left (0):
{
    float dtemp = ( k_over_rho_c *
                    ( left - 2.*PPTemps[0] + PPTemps[1] ) / ( DX*DX ) ) * DT;
    NextTemps[0] = PPTemps[0] + dtemp;
}

// all the nodes in the middle:
for( int i = 1; i < PPSIZE-1; i++ )
{
    float dtemp = ( k_over_rho_c *
                    ( PPTemps[i-1] - 2.*PPTemps[ i ] + PPTemps[i+1] ) / ( DX*DX ) ) * DT;
    NextTemps[i] = PPTemps[i] + dtemp;
}

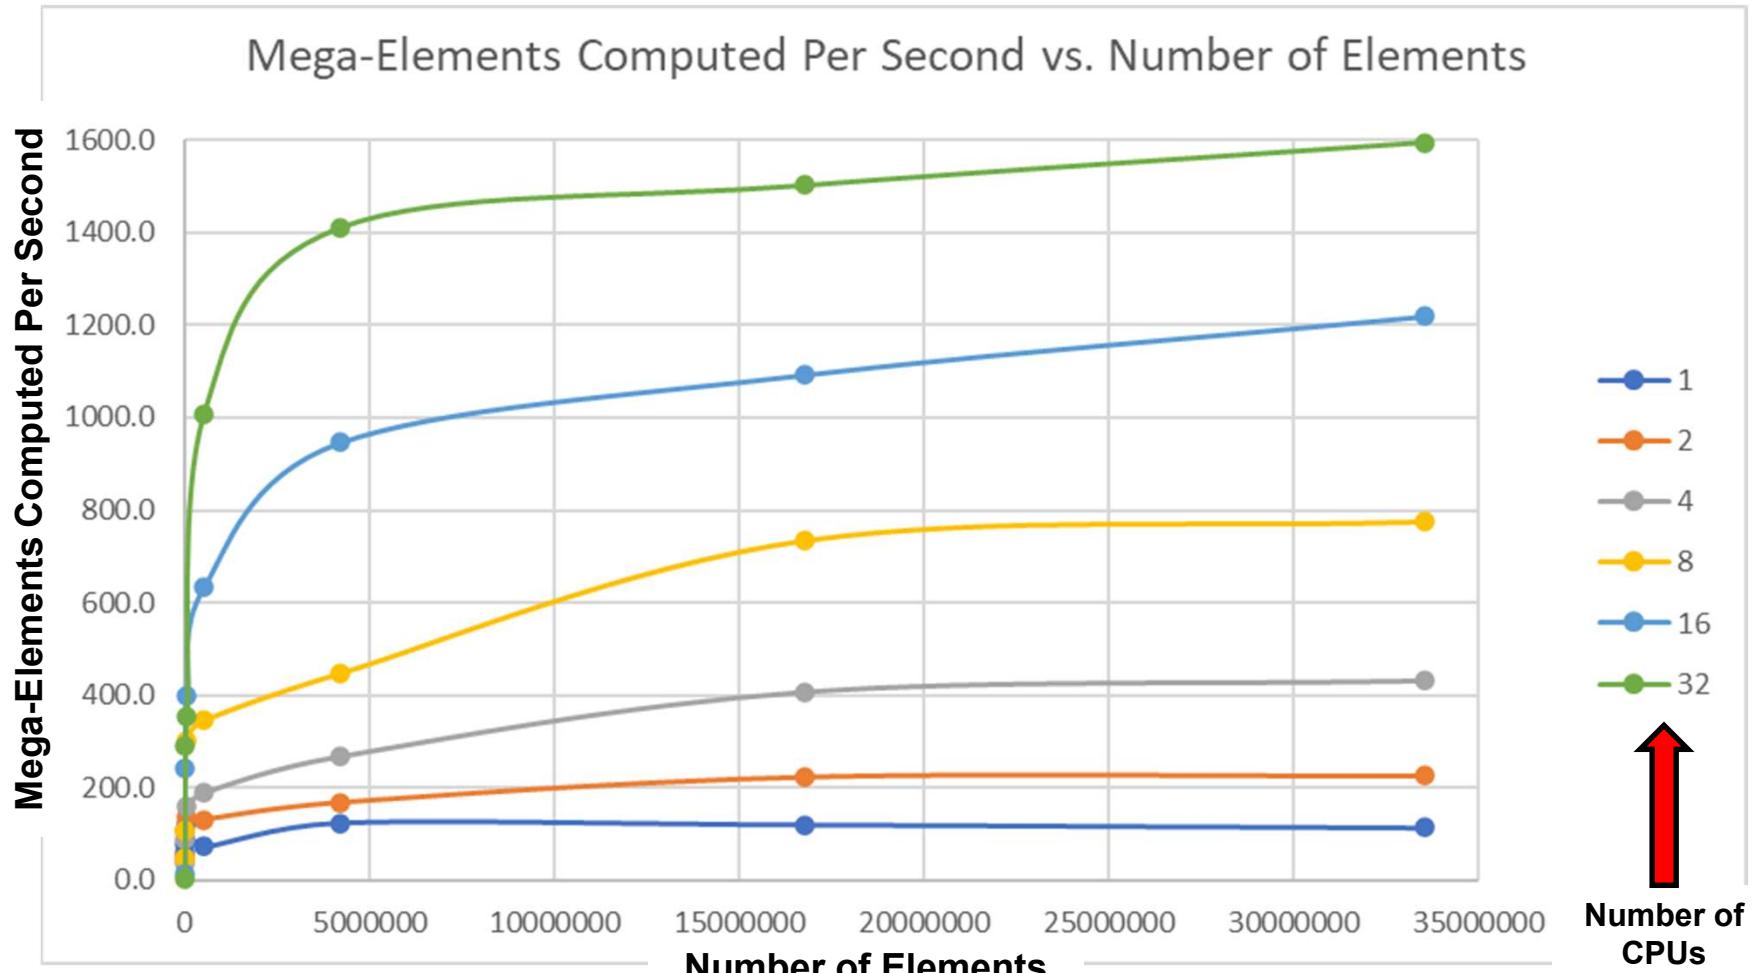
// last element on the right (PPSIZE-1):
{
    float dtemp = ( k_over_rho_c *
                    ( PPTemps[PPSIZE-2] - 2.*PPTemps[PPSIZE-1] + right ) / ( DX*DX ) ) * DT;
    NextTemps[PPSIZE-1] = PPTemps[PPSIZE-1] + dtemp;
}
```

DoOneTimeStep, IV

```
// update the local dataset:  
  
for( int i = 0; i < PPSize; i++ )  
{  
    PPTemps[ i ] = NextTemps[ i ];  
}  
}
```

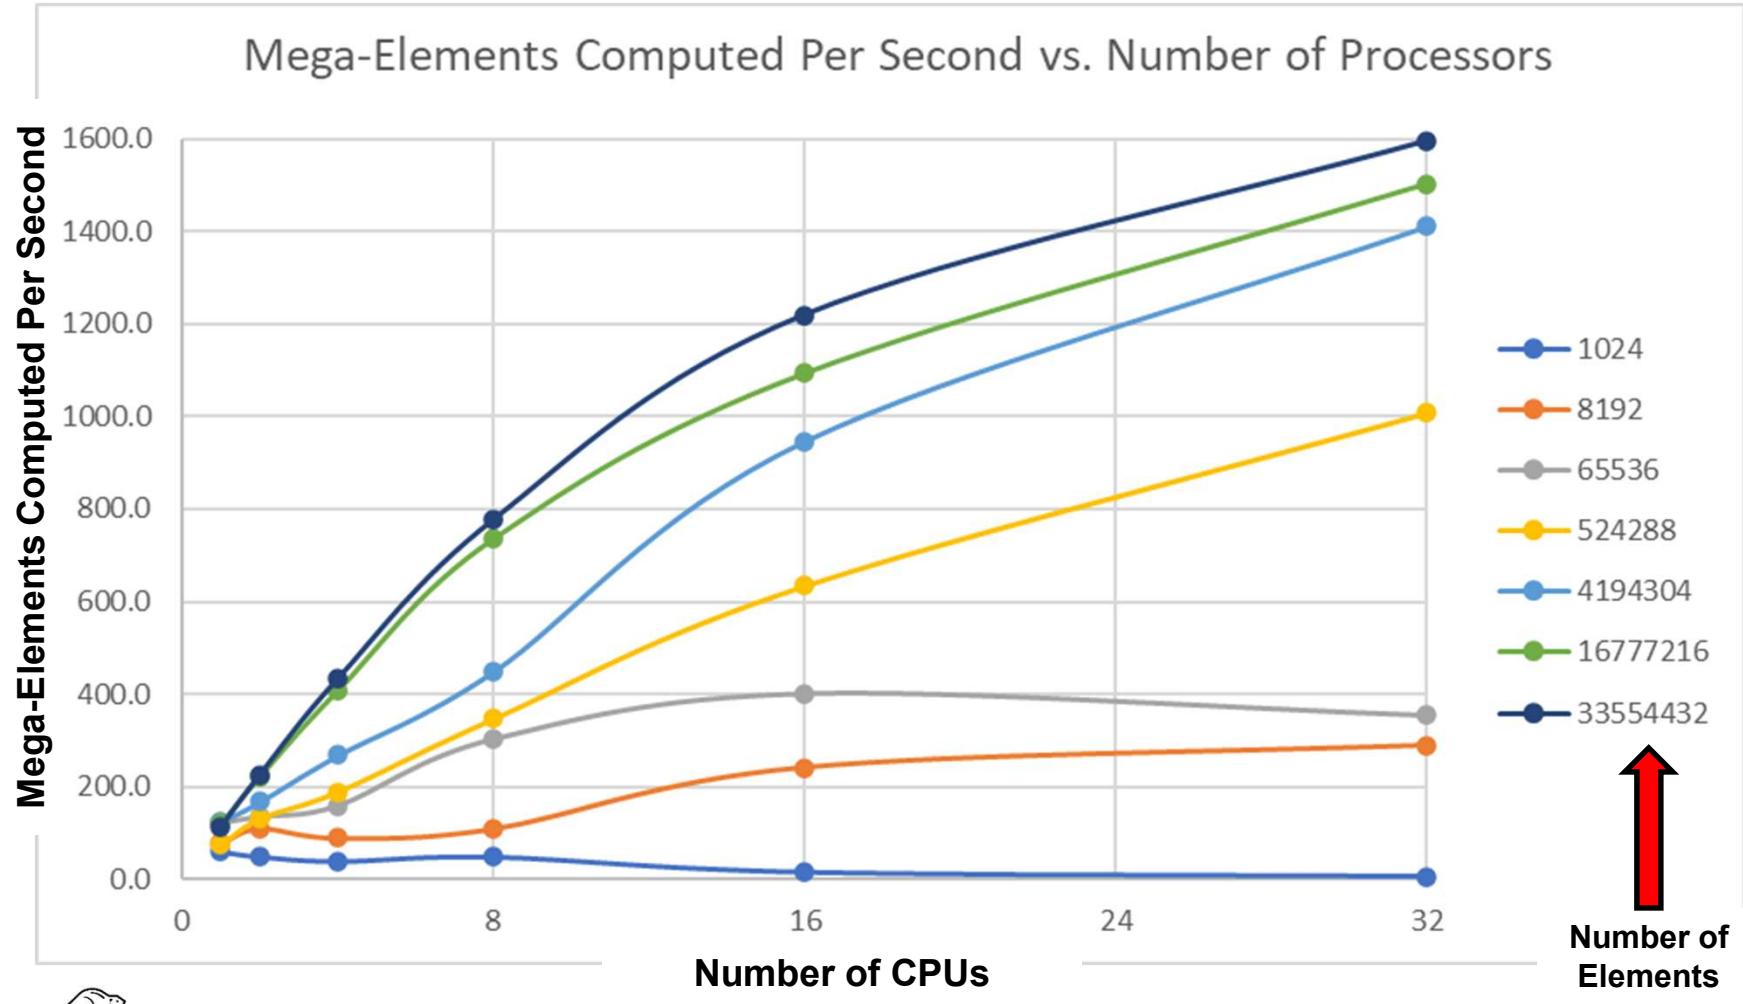
MPI Performance

42



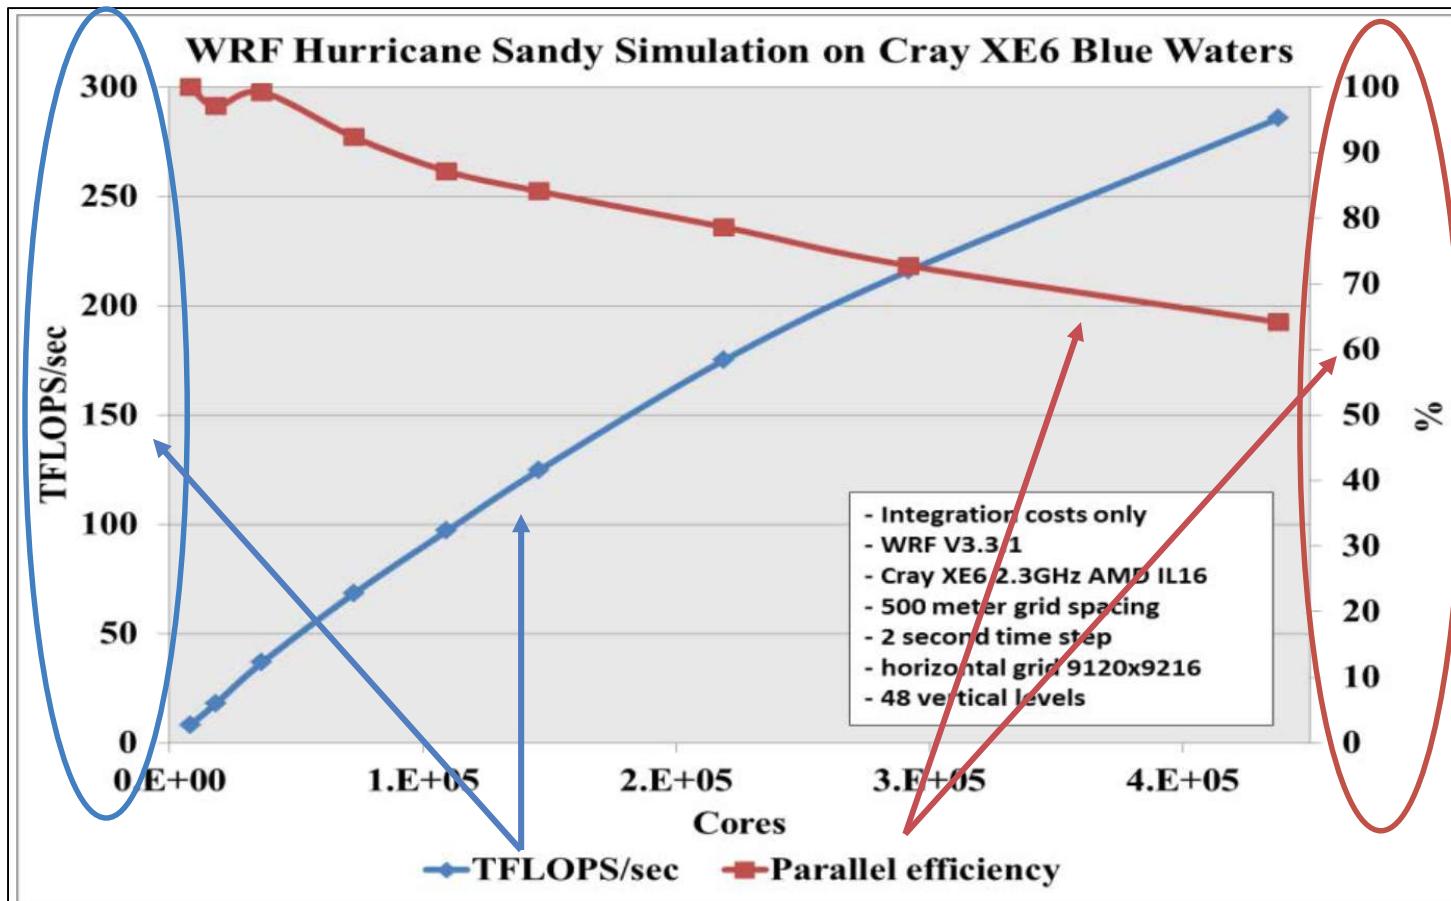
MPI Performance

43



Using MPI and OpenMP on 13,680 nodes (437,760 cores) of the Cray XE6 at NCSA at the University of Illinois

44



From: Peter Johnsen, Mark Straka, Melvyn Shapiro, Alan Norton, Thomas Galarneau,
Petascale WRF Simulation of Hurricane Sandy.



Oregon State
University
Computer Graphics

MPI Reduction

```
MPI_Reduce( partialResult, globalResult, count, type, operator, dst, MPI_COMM_WORLD );
```

Where the partial result is stored on each CPU

Place to store the full result on the dst CPU

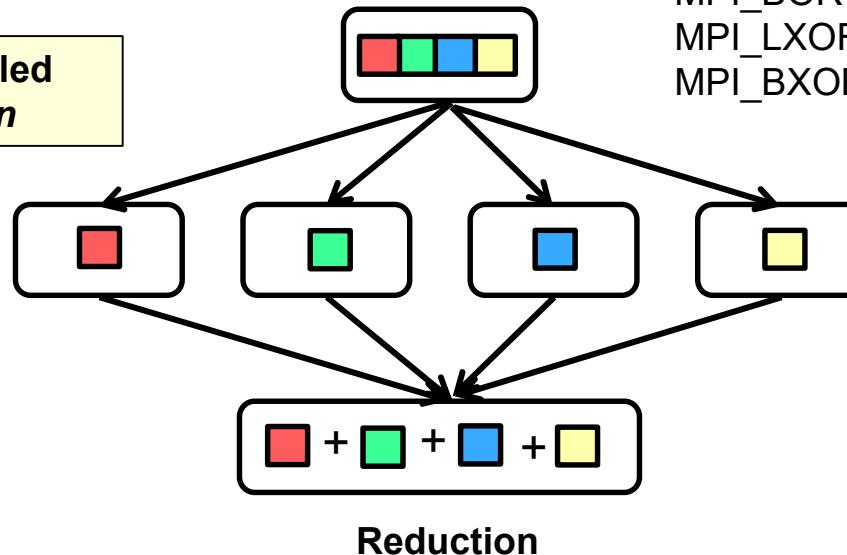
Number of elements in the partial result

MPI_CHAR
MPI_INT
MPI_LONG
MPI_FLOAT
MPI_DOUBLE
• • •

MPI_MIN
MPI_MAX
MPI_SUM
MPI_PROD
MPI_MINLOC
MPI_MAXLOC
MPI_LAND
MPI_BAND
MPI_LOR
MPI_BOR
MPI_LXOR
MPI_BXOR

Who is given the final answer

This really should be called
Scatter/Gather/Reduction



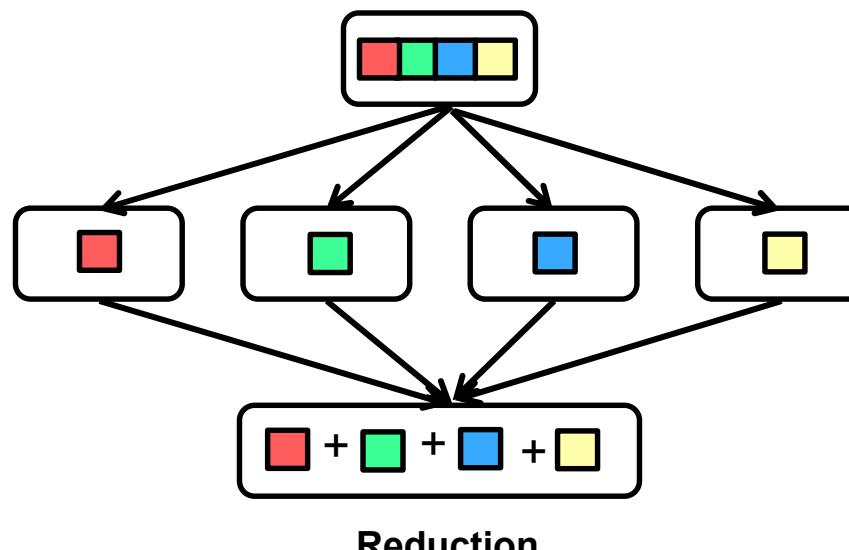
MPI Reduction Example

```
// gratuitous use of a reduce -- average all the temperatures:

float partialSum = 0.;
for( int i = 0; i < PPSIZE; i++ )
    partialSum += PPTemps[ i ];

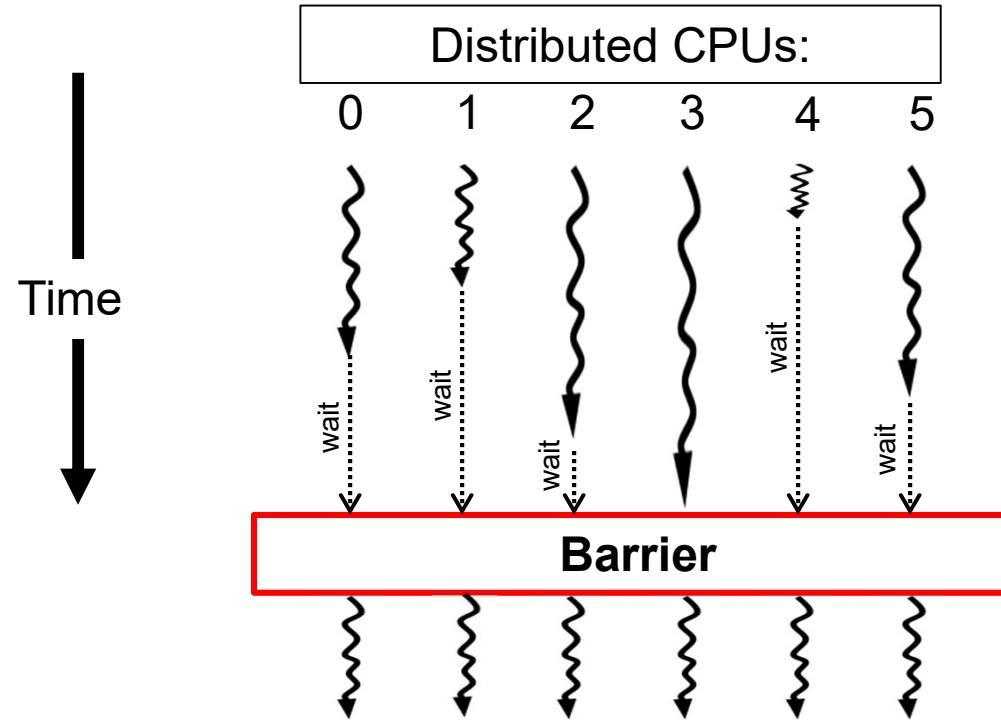
float globalSum = 0.;
MPI_Reduce( &partialSum, &globalSum, 1, MPI_FLOAT, MPI_SUM, THEBOSS, MPI_COMM_WORLD );

if( me == THEBOSS )
    fprintf( stderr, "Average temperature = %f\n", globalSum/(float)NUMELEMENTS );
```



MPI Barriers

```
MPI_Barrier( MPI_COMM_WORLD );
```



All CPUs must execute the call to `MPI_Barrier()` before any of the CPUs can move past it. That is, each CPU's `MPI_Barrier()` blocks until all CPUs execute a call to `MPI_Barrier()`.

MPI Derived Types

Idea: In addition to types MPI_INT, MPI_FLOAT, etc., allow the creation of new MPI types so that you can transmit an “array of structures”.

Reason: There is significant overhead with each transmission. Better to send one entire array of structures instead of sending several arrays separately.

```
MPI_Type_create_struct( count, blocklengths, displacements, types, datatype );
```

```
struct point
{
    int pointSize;
    float x, y, z;
};
```

```
MPI_Datatype MPI_POINT;
int blocklengths[ ] = { 1, 1, 1, 1 };
int displacements[ ] = { 0, 4, 8, 12 };
MPI_type types[ ] = { MPI_INT, MPI_FLOAT, MPI_FLOAT, MPI_FLOAT };

MPI_Type_create_struct( 4, blocklengths, displacements, types, &MPI_POINT )
```



MPI Timing

```
double MPI_Wtick( );
```

Returns the resolution of the clock, in seconds.

```
double MPI_Wtime( );
```

Returns the time, in seconds, since “some time in the past”.

Warning: the clocks on the different CPUs are not guaranteed to be synchronized!



Oregon State
University
Computer Graphics

MPI Status-Checking

Some MPI calls have a **&status** in their argument list.

The **status** argument is declared to be of type **MPI_Status**, which is actually a struct:

```
typedef struct _MPI_Status
{
    int MPI_SOURCE;
    int MPI_TAG;
    int MPI_ERROR;
} MPI_Status;
```

- **MPI_SOURCE** is the rank of the node who sent this
- **MPI_TAG** is the tag used during the send
- **MPI_ERROR** is the error number that occurred

Example:

```
MPI_Status status;
MPI_Recv( myData, MYDATA_SIZE, MPI_CHAR, THEBOSS, MPI_ANY_TAG, MPI_COMM_WORLD,
&status );

fprintf( stderr, "Rank = %d, Tag = %d, Error Code = %d\n",
        status(MPI_SOURCE, status.MPI_TAG, status.MPI_ERROR );
```



Oregon State
University
Computer Graphics

MPI Error Codes

<code>MPI_SUCCESS</code>	No error	
<code>MPI_ERR_BUFFER</code>	Invalid buffer pointer	
<code>MPI_ERR_COUNT</code>	Invalid count argument	
<code>MPI_ERR_TYPE</code>	Invalid datatype argument	
<code>MPI_ERR_TAG</code>	Invalid tag argument	
<code>MPI_ERR_COMM</code>	Invalid communicator	
<code>MPI_ERR_RANK</code>	Invalid rank	
<code>MPI_ERR_REQUEST</code>	Invalid request (handle)	
<code>MPI_ERR_ROOT</code>	Invalid root	
<code>MPI_ERR_GROUP</code>	Invalid group	
<code>MPI_ERR_OP</code>	Invalid operation	
<code>MPI_ERR_TOPOLOGY</code>	Invalid topology	
<code>MPI_ERR_DIMS</code>	Invalid dimension argument	
<code>MPI_ERR_ARG</code>	Invalid argument of some other kind	
<code>MPI_ERR_UNKNOWN</code>	Unknown error	
<code>MPI_ERR_TRUNCATE</code>	Message truncated on receive	
<code>MPI_ERR_OTHER</code>	Known error not in this list	
<code>MPI_ERR_INTERNAL</code>	Internal MPI (implementation) error	
<code>MPI_ERR_IN_STATUS</code>	Error code is in status	
<code>MPI_ERR_PENDING</code>	Pending request	
<code>MPI_ERR_FILE</code>		Invalid file handle
<code>MPI_ERR_NOT_SAME</code>		Collective argument not identical on all processes, or collective routines called in a different order by different processes
<code>MPI_ERR_AMODE</code>		Error related to the amode passed to <code>MPI_FILE_OPEN</code>
<code>MPI_ERR_UNSUPPORTED_DATAREP</code>		Unsupported datarep passed to <code>MPI_FILE_SET_VIEW</code>
<code>MPI_ERR_UNSUPPORTED_OPERATION</code>		Unsupported operation, such as seeking on a file which supports sequential access only
<code>MPI_ERR_NO SUCH FILE</code>		File does not exist
<code>MPI_ERR_FILE_EXISTS</code>		File exists
<code>MPI_ERR_BAD_FILE</code>		Invalid file name (e.g., path name too long)
<code>MPI_ERR_ACCESS</code>		Permission denied
<code>MPI_ERR_NO_SPACE</code>		Not enough space
<code>MPI_ERR_QUOTA</code>		Quota exceeded
<code>MPI_ERR_READ_ONLY</code>		Read-only file or file system
<code>MPI_ERR_FILE_IN_USE</code>		File operation could not be completed, as the file is currently open by some process
<code>MPI_ERR_DUP_DATAREP</code>		Conversion functions could not be registered because a data representation identifier that was already defined was passed to <code>MPI_REGISTER_DATAREP</code>
 		An error occurred in a user supplied data conversion function.
<code>Oregon</code>	<code>MPI_ERR_CONVERSION</code>	Other I/O error
<code>Univer</code>	<code>MPI_ERR_IO</code>	Last error code
<code>Computer</code>	<code>MPI_ERR_LASTCODE</code>	

