

Performing Reductions in OpenCL (and CUDA)



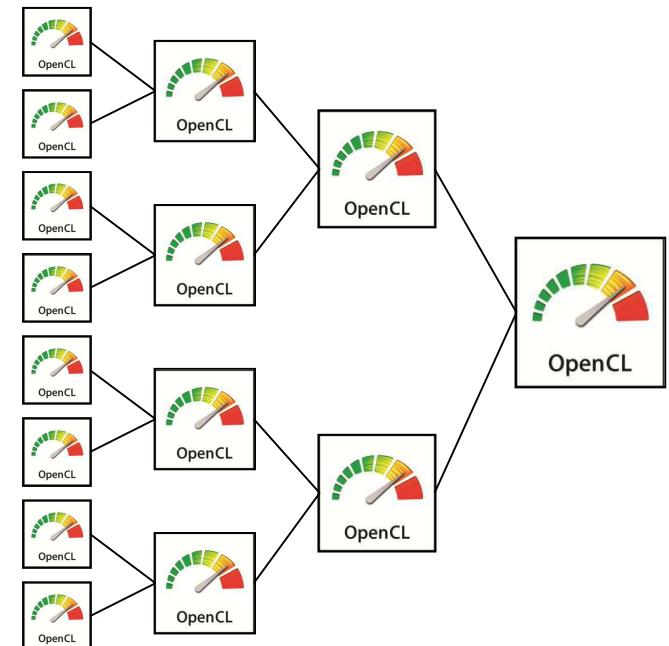
Oregon State
University
Mike Bailey

mjb@cs.oregonstate.edu

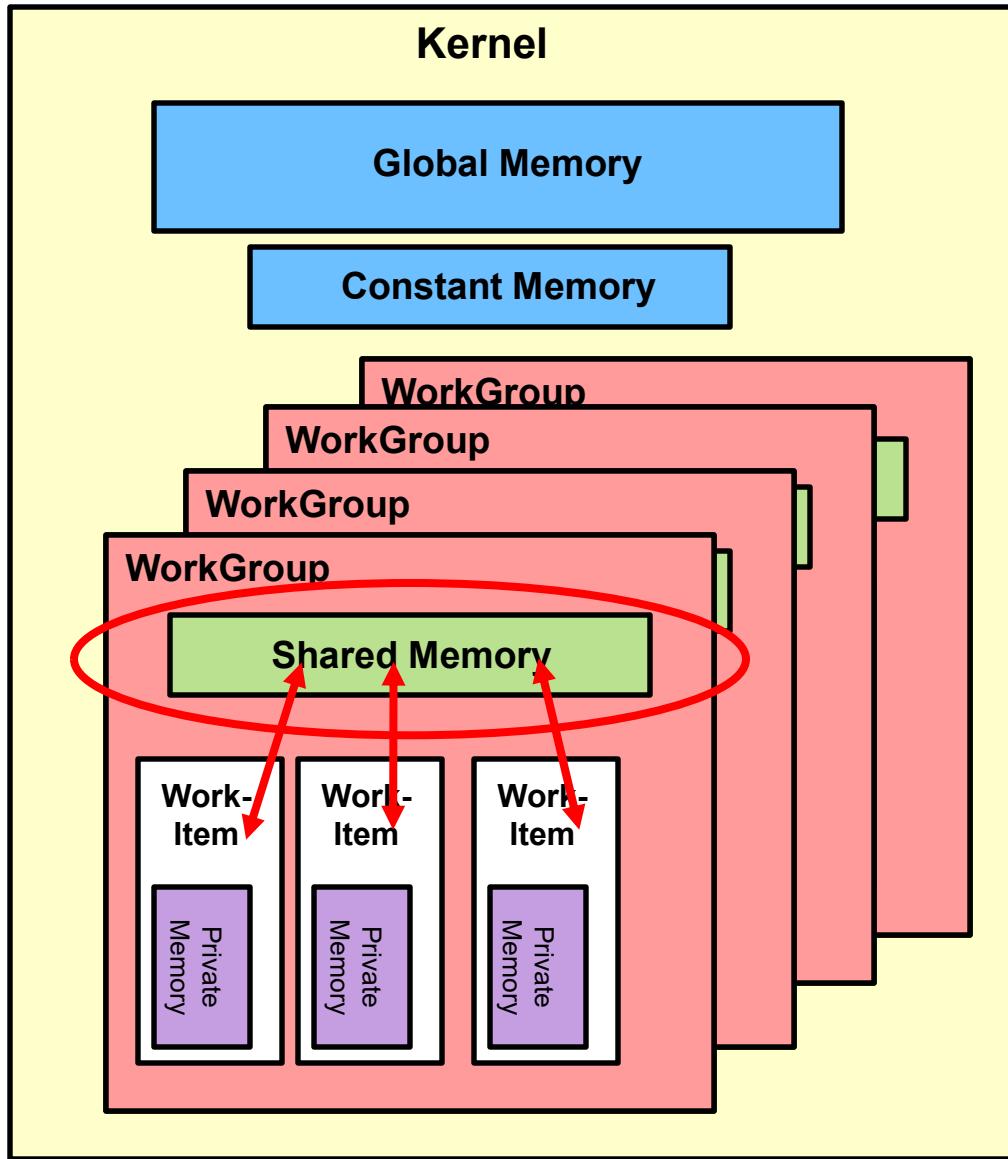


This work is licensed under a [Creative Commons](#)
[Attribution-NonCommercial-NoDerivatives 4.0](#)
[International License](#)

Oregon State
University
Computer Graphics



Recall the OpenCL Memory Model



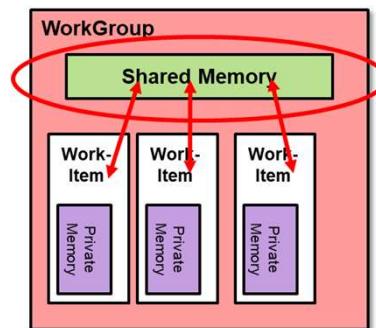
Here's the Problem We're Trying to Solve

Like the *first.cpp* demo program, we are piecewise multiplying two arrays. Unlike the first demo program, we want to then add up all the products and return the sum.

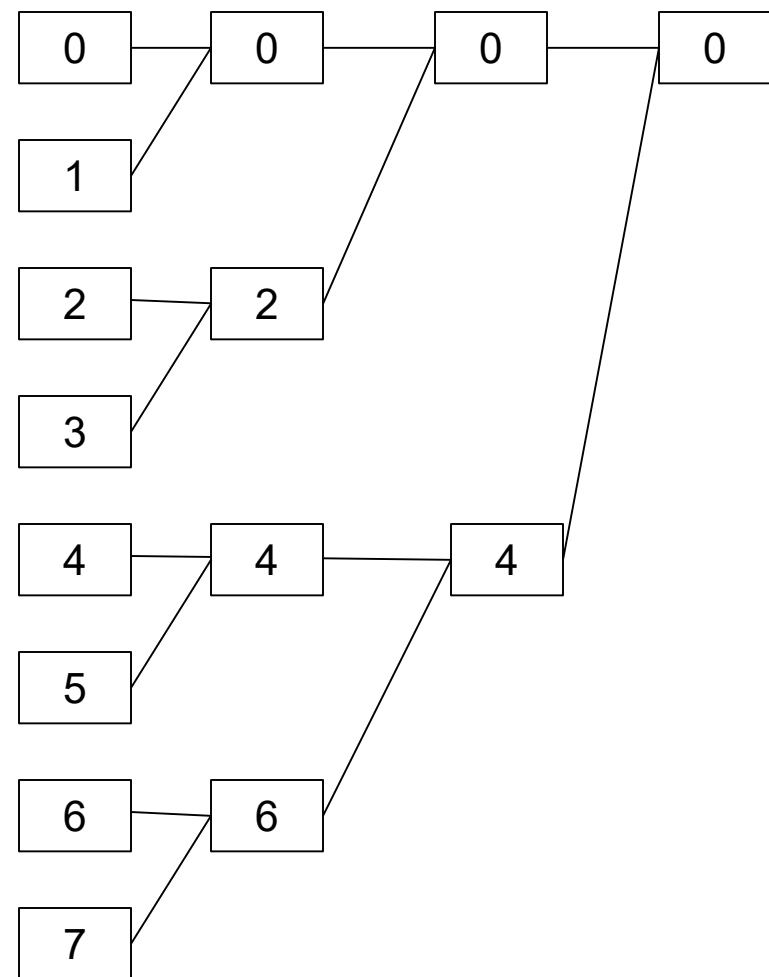
$$\begin{aligned} A * B &\rightarrow \text{prods} \\ \sum \text{prods} &\rightarrow C \end{aligned}$$

After the array multiplication, we want each work-group to sum the products within that work-group, then return them to the host in an array for final summing.

To do this, we will not put the products into a large global device array, but into a **prods[]** array that is shared within its work-group.

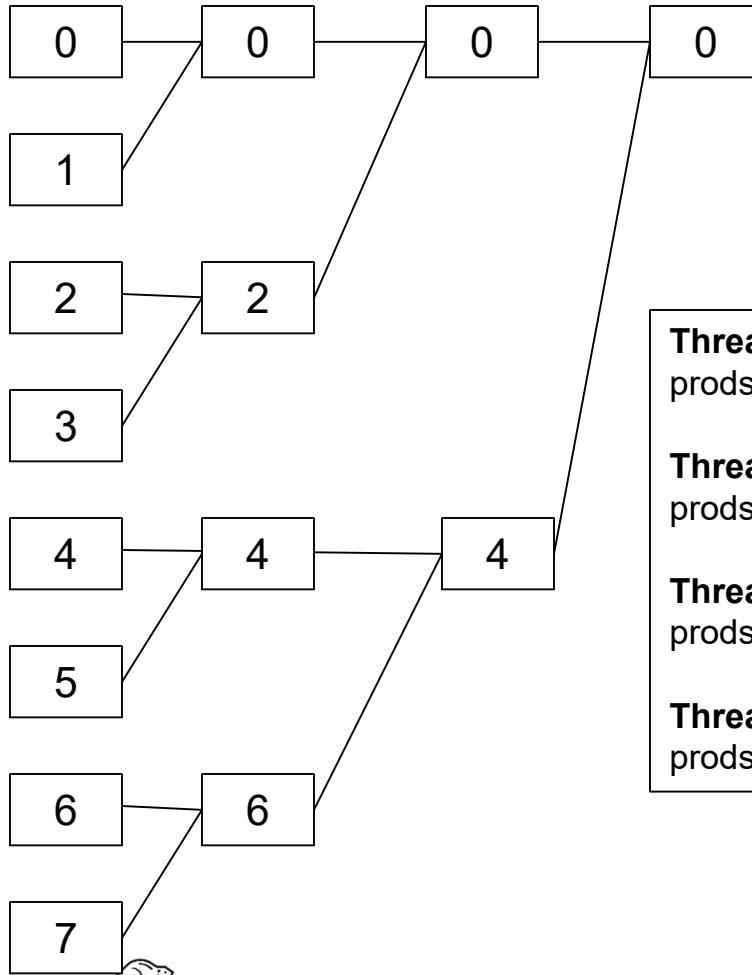


`numItems = 8;`



Reduction Takes Place in a Single Work-Group

numItems = 8;



If we had 8 work-items in a work-group, we would like the threads in each work-group to execute the following instructions . . .

Thread #0: prods[0] += prods[1];	Thread #0: prods[0] += prods[2];	Thread #0: prods[0] += prods[4];
Thread #2: prods[2] += prods[3];		
Thread #4: prods[4] += prods[5];	Thread #4: prods[4] += prods[6];	
Thread #6: prods[6] += prods[7];		

. . . but in a more general way than writing them all out by hand.



Here's What You Would Change in your Host Program

```
#define NUM_WORKGROUPS ( NUM_ELEMENTS / LOCAL_SIZE )

// global variables:
float hA[ NUM_ELEMENTS ];
float hB[ NUM_ELEMENTS ];
float hC[NUM_WORKGROUPS ];

• • •
size_t abSize = NUM_ELEMENTS      * sizeof(float);
size_t cSize = NUM_WORKGROUPS    * sizeof(float);

• • •
cl_mem dA = clCreateBuffer( context, CL_MEM_READ_ONLY, abSize, NULL, &status );
cl_mem dB = clCreateBuffer( context, CL_MEM_READ_ONLY, abSize, NULL, &status );
cl_mem dC = clCreateBuffer( context, CL_MEM_WRITE_ONLY, cSize, NULL, &status );

• • •
status = clEnqueueWriteBuffer( cmdQueue, dA, CL_FALSE, 0, abSize, hA, 0, NULL, NULL );
status = clEnqueueWriteBuffer( cmdQueue, dB, CL_FALSE, 0, abSize, hB, 0, NULL, NULL );
• • •
cl_kernel kernel = clCreateKernel( program, "ArrayMultReduce", &status );
• • •

status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );
status = clSetKernelArg( kernel, 2, LOCAL_SIZE * sizeof(float), NULL );  

        // local "prods" array is dimensioned the size of each work-group
status = clSetKernelArg( kernel, 3, sizeof(cl_mem), &dC );
```

$$A * B \rightarrow \text{prods}$$

$$\sum \text{prods} \rightarrow C$$

"cl_mem" is a GPU buffer
memory address

This NULL is how you tell
OpenCL that this is a *local*/
(shared) array, not a global array



The Arguments to the Kernel

```

status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );
status = clSetKernelArg( kernel, 2, LOCAL_SIZE * sizeof(float), NULL );
                                // local "prods" array – one per work-item
status = clSetKernelArg( kernel, 3, sizeof(cl_mem), &dC );

```

kernel void
ArrayMultReduce(global const float *dA, global const float *dB, **local** float *prods, global float *dC)
{
 int gid = get_global_id(0); // 0 .. total_array_size-1
 int numItems = get_local_size(0); // # work-items per work-group
 int tnum = get_local_id(0); // thread (i.e., work-item) number in this work-group
 // 0 .. numItems-1
 int wgNum = get_group_id(0); // which work-group number this is in

 prods[tnum] = dA[gid] * dB[gid]; // multiply the two arrays together

$A * B \rightarrow \text{prods}$

// now add them up – come up with one sum per work-group
// it is a big performance benefit to do it here while "prods" is still available – and is local
// it would be a performance hit to pass "prods" back to the host then bring it back to the device for reduction



Reduction Takes Place Within a Single Work-Group

Each work-item is run by a single thread

Thread #0:
prods[0] += prods[1];

Thread #2:
prods[2] += prods[3];

Thread #4:
prods[4] += prods[5];

Thread #6:
prods[6] += prods[7];

offset = 1;
mask = 1;

Thread #0:
prods[0] += prods[2];

Thread #4:
prods[4] += prods[6];

offset = 2;
mask = 3;

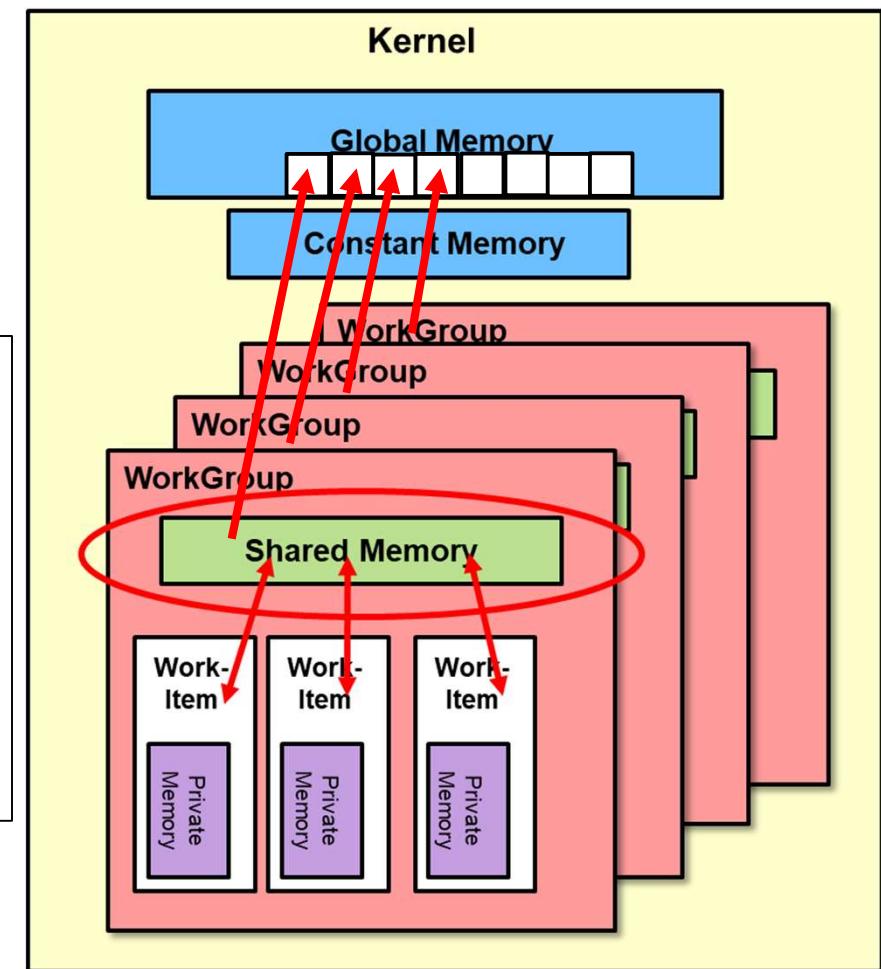
Thread #0:
prods[0] += prods[4];

offset = 4;
mask = 7;

A work-group consisting of *numItems* work-items can be reduced to a sum in $\log_2(\text{numItems})$ steps. In this example, *numItems*=8.

The reduction begins with the individual products in prods[0] .. prods[7].

The final sum will end up in prods[0], which will then be copied into dC[wgNum].



A Review of Bitmasks

Remember *Truth Tables*?

F	& F
<hr/>	
= F	

F	& T
<hr/>	
= F	

T	& F
<hr/>	
= F	

T	& T
<hr/>	
= T	

Or, with Bits:

0	& 0
<hr/>	
= 0	

0	& 1
<hr/>	
= 0	

1	& 0
<hr/>	
= 0	

1	& 1
<hr/>	
= 1	

Or, with Multiple Bits:

000	& 011
<hr/>	
= 000	

001	& 011
<hr/>	
= 001	

010	& 011
<hr/>	
= 010	

011	& 011
<hr/>	
= 011	

100	& 011
<hr/>	
= 000	

101	& 011
<hr/>	
= 001	



If it's been a long time since you have looked at bitmask operators (or never!), here is a good review reference:
https://en.wikipedia.org/wiki/Bitwise_operations_in_C

Reduction Takes Place in a Single *Work-Group*

Each *work-item* is run by a single thread

numItems = 8;

Thread #0: prods[0] += prods[1];	Thread #0: prods[0] += prods[2];	Thread #0: prods[0] += prods[4];
Thread #2: prods[2] += prods[3];		offset = 4; mask = 7;
Thread #4: prods[4] += prods[5];	Thread #4: prods[4] += prods[6];	
Thread #6: prods[6] += prods[7];		offset = 2; mask = 3;
offset = 1; mask = 1;		



Reduction Takes Place in a Single Work-Group

Each *work-item* is run by a single thread

Thread #0: prods[0] += prods[1];	Thread #0: prods[0] += prods[2];	Thread #0: prods[0] += prods[4]; offset = 4; mask = 7;
Thread #2: prods[2] += prods[3];		
Thread #4: prods[4] += prods[5];	Thread #4: prods[4] += prods[6]; offset = 2; mask = 3;	
Thread #6: prods[6] += prods[7];		
offset = 1; mask = 1;		

```

kernel void
ArrayMultReduce( ... )
    int gid      = get_global_id( 0 );
    int numItems = get_local_size( 0 );
    int tnum     = get_local_id( 0 ); // thread number
    int wgNum    = get_group_id( 0 ); // work-group number

    // all threads execute this code simultaneously:
    prods[ tnum ] = dA[ gid ] * dB[ gid ];
    for( int offset = 1; offset < numItems; offset *= 2 )
    {
        int mask = 2*offset - 1;
        barrier( CLK_LOCAL_MEM_FENCE ); // wait for all threads to get here
        if( ( tnum & mask ) == 0 )           // bit-by-bit and'ing tells us which
                                              // threads need to do work now
        {
            prods[ tnum ] += prods[ tnum + offset ];
        }
    }

    barrier( CLK_LOCAL_MEM_FENCE );
    if( tnum == 0 )
        dC[ wgNum ] = prods[ 0 ];
}

```

Anding bits → $\sum \text{prods} \rightarrow C$

And, Finally, in your Host Program

```
Wait( cmdQueue );
double time0 = omp_get_wtime( );

status = clEnqueueNDRangeKernel( cmdQueue, kernel, 1, NULL, globalWorkSize, localWorkSize,
                                0, NULL, NULL );
PrintCLError( status, "clEnqueueNDRangeKernel failed: " );

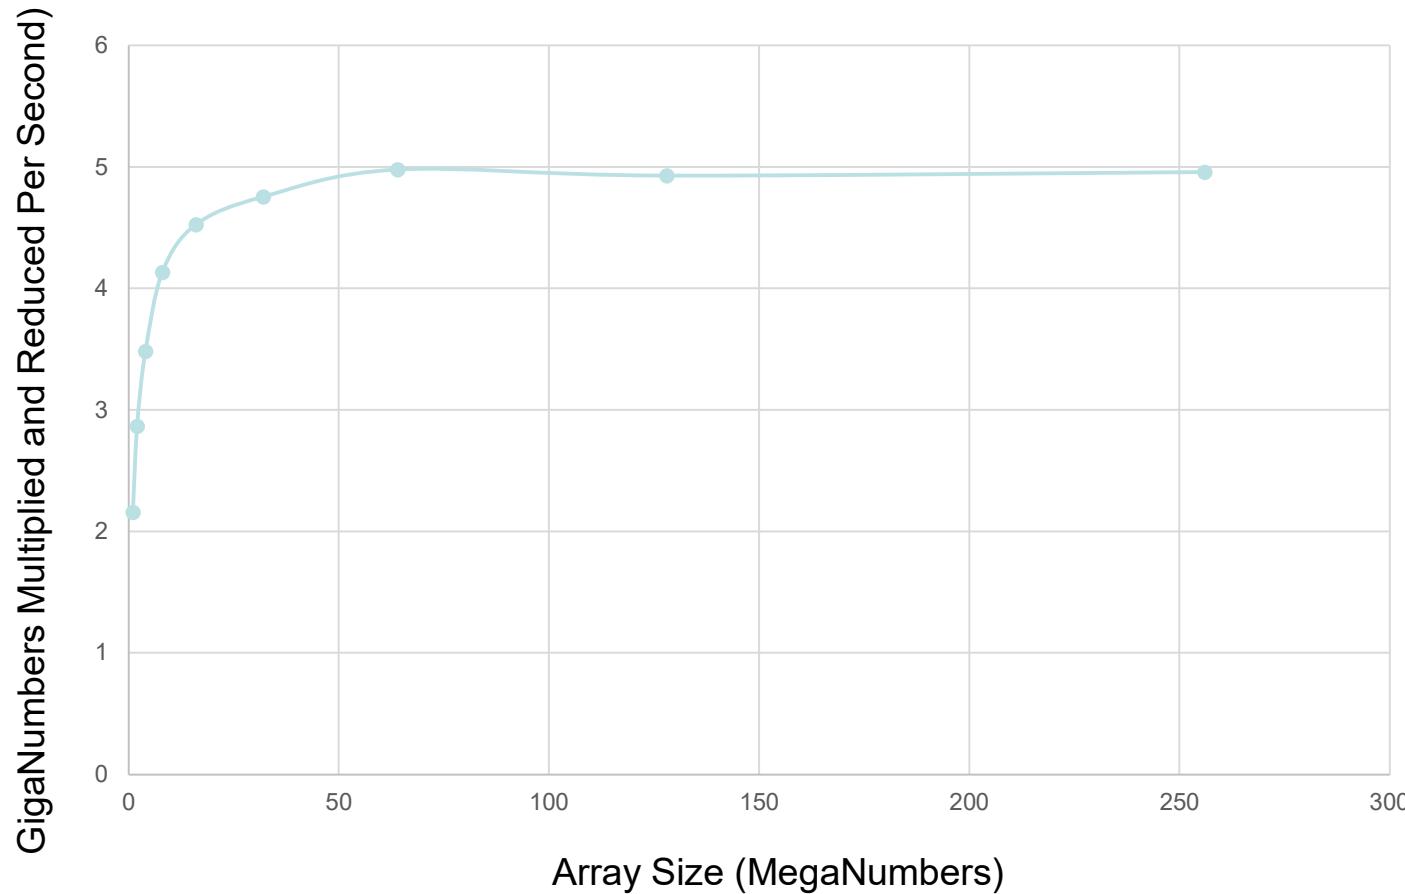
Wait( cmdQueue );
double time1 = omp_get_wtime( );

status = clEnqueueReadBuffer( cmdQueue, dC, CL_TRUE, 0, NUM_WORKGROUPS*sizeof(float), hC,
                             0, NULL, NULL );
PrintCLError( status, "clEnqueueReadBuffer failed: " );
Wait( cmdQueue );

float sum = 0.;
for( int i = 0; i < numWorkgroups; i++ )
{
    sum += hC[ i ];
}
```

Reduction Performance

Work-Group Size = 32



Oregon State
University
Computer Graphics



mjb – May 15, 2025