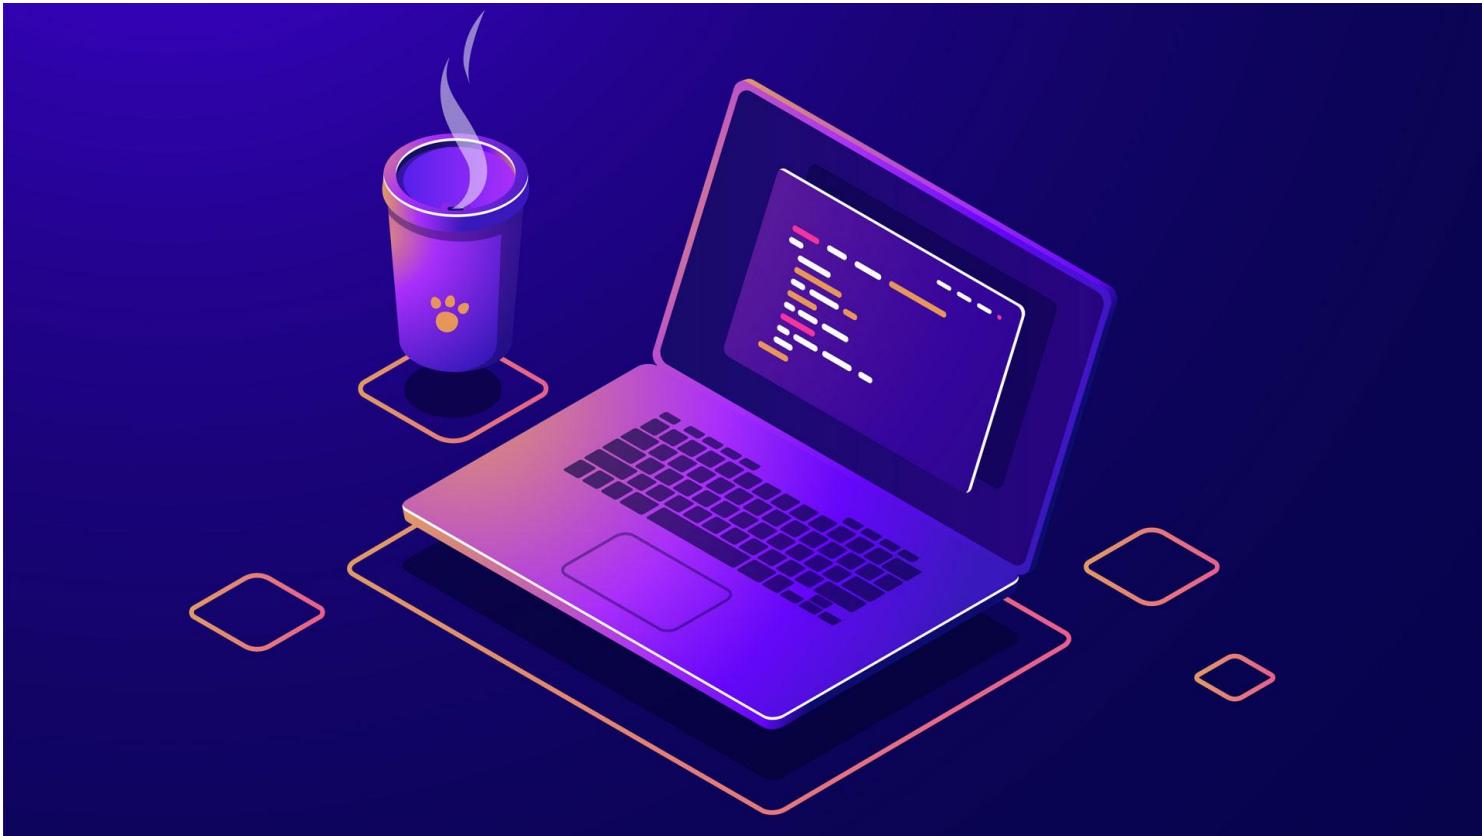


Advanced Granularity



Advanced Granularity



Given a sequential program, the number of tasks that one can generate and the size of the tasks (what is called **granularity**) are related one to the other.

- ▶ Fine-grained tasks vs. coarse-grained tasks
- ▶ The parallelism increases as the decomposition becomes finer in granularity (small tasks) and vice versa

Example 1: matrix-vector product




Fine-grained Decomposition



Example: matrix-vector product (n by n matrix):

- ▶ A task could be each individual \times and $+$ in the dot product that contributes to the computation of an element of y
 $(y[i] = y[i] + A[i][j] * b[j])$




- ▶ A task could also be each complete dot product to compute an element of y ($y[i] = y[i] + \sum_{j=1}^{j=n} (A[i][j] * b[j])$)

Coarse-grained Decomposition



- ▶ A task could be in charge of computing a number of consecutive elements of y (e.g. three elements)



- ▶ A task could be in charge of computing the whole vector y

So...



- ▶ It would appear that the parallel time can be made arbitrarily small by making the decomposition finer in granularity but...
 - ▶ Inherent bound on how fine the granularity of a computation can be
 - ▶ e.g. *matrix-vector multiply*: (n^2) concurrent tasks.
 - ▶ Tradeoff between the granularity of a decomposition and associated overheads (sources of overhead: creation of tasks, task synchronization, exchange of data between tasks, ...)
 - ▶ The granularity may determine performance bounds

Example 2: stencil computation using Jacobi solver

Stencil algorithm that computes each element of matrix `utmp` using 4 neighbor elements of matrix `u`, both matrices with $n \times n$ elements

```
void compute(int n, double *u, double *utmp) {
    int i, j;
    double tmp;

    for (i = 1; i < n-1; i++) {
        for (j = 1; j < n-1; j++) {
            tmp = u[n*(i+1) + j] + u[n*(i-1) + j] + // elements u[i+1][j] and u[i-1][j]
                  u[n*i + (j+1)] + u[n*i + (j-1)] - // elements u[i][j+1] and u[i][j-1]
                  4 * u[n*i + j];                      // element u[i][j]
            utmp[n*i + j] = tmp / 4;                // element utmp[i][j]
        }
    }
}
```

Example 2: stencil computation using Jacobi solver

What tasks can be? Assume: 1) the innermost loop body takes t_{body} time units; and 2) n is very large, so that $n - 2 \simeq n$

| Task is ... (granularity) | Num. tasks | Task cost | T_1 | T_∞ | Parallelism |
|---|------------------------|----------------------------|----------------------|----------------------------|------------------------|
| All iterations of i and j loops | 1 | $n^2 \cdot t_{body}$ | $n^2 \cdot t_{body}$ | $n^2 \cdot t_{body}$ | 1 |
| Each iteration of i loop | n | $n \cdot t_{body}$ | $n^2 \cdot t_{body}$ | $n \cdot t_{body}$ | n |
| Each iteration of j loop | n^2 | t_{body} | $n^2 \cdot t_{body}$ | t_{body} | n^2 |
| r consecutive iterations of i loop | $n \div r$ | $n \cdot r \cdot t_{body}$ | $n^2 \cdot t_{body}$ | $n \cdot r \cdot t_{body}$ | $n \div r$ |
| c consecutive iterations of j loop | $n^2 \div c$ | $c \cdot t_{body}$ | $n^2 \cdot t_{body}$ | $c \cdot t_{body}$ | $n^2 \div c$ |
| A block of $r \times c$ iterations of i and j, respectively | $n^2 \div (r \cdot c)$ | $r \cdot c \cdot t_{body}$ | $n^2 \cdot t_{body}$ | $r \cdot c \cdot t_{body}$ | $n^2 \div (r \cdot c)$ |

Finer grain task decomposition → higher parallelism, but ...

Example 2: stencil computation using Jacobi solver

... what if each task creation takes t_{create} ?

| Task is ... (granularity) | Num. tasks | Task cost | Task creation ovh |
|---|------------------------|----------------------------|---|
| All iterations of i and j loops | 1 | $n^2 \cdot t_{body}$ | t_{create} |
| Each iteration of i loop | n | $n \cdot t_{body}$ | $n \cdot t_{create}$ |
| Each iteration of j loop | n^2 | t_{body} | $n^2 \cdot t_{create}$ |
| r consecutive iterations of i loop | $n \div r$ | $n \cdot r \cdot t_{body}$ | $(n \div r) \cdot t_{create}$ |
| c consecutive iterations of j loop | $n^2 \div c$ | $c \cdot t_{body}$ | $(n^2 \div c) \cdot t_{create}$ |
| A block of $r \times c$ iterations of i and j, respectively | $n^2 \div (r \cdot c)$ | $r \cdot c \cdot t_{body}$ | $(n^2 \div (r \cdot c)) \cdot t_{create}$ |

Trade-off between task granularity and task creation overhead

Instructor Social Media

Youtube: Lucas Science



Instagram: lucaasbazilio



Twitter: lucasebazilio

